

Using Machine Learning Tools PG

Week 6 – Logistic Regression &
Support Vector Machines

COMP SCI 7317

Trimester 2, 2024



THE UNIVERSITY
of ADELAIDE

150 YEARS

From last week... Training models

1. Training Models

- Cost/error/loss function
- Gradient and stochastic gradient descent continued
- Learning rate continued
- Stopping criteria
- Training curve/learning curve

2. Regularisation

- L2 regularisation (ridge/Tikhonov)
- L1 regularisation (Lasso)
- Elastic net

This week

1. Regularisation Continued

- L1, L2, Elastic Net

2. Logistic Regression

- Logistic & Logit Regression
- Training a logistic model
- Softmax

3. Support Vector Machines (SVM)

- Linear SVM - Hard Margin
- Linear SVM - Soft Margin
- Kernel trick



THE UNIVERSITY
of ADELAIDE

150 YEARS



Regularisation Continued

Regularisation

Regularisation: a method that adds a term to the cost function to prevent overfitting. Controlled by an adjustable weight α .

$$Cost = data_term + \alpha * regularisation_term$$

Purpose:

- Help avoid overfitting by penalising large parameter values.
- Encourages 'smooth' outputs.
- Add desired properties to the model (a-priori knowledge).
- Balance between fitting data and simplicity (multi-objective optimisation).

L2 (Ridge/Tikhonov) Regularisation

Effect: Penalises large model parameters to prevent overfitting.

Cost **Data term** **Regularisation term**



$$J(\boldsymbol{\theta}) = \text{MSE}(\boldsymbol{\theta}) + \alpha \frac{1}{2} \sum_{i=1}^n \theta_i^2$$

L2 Norm

α : Regularisation strength or penalty term

$\boldsymbol{\theta}$: Model parameters, in this case excluding θ_0

- Scaling of data important for setting α
- Scikit-learn: `penalty` parameter “l2”



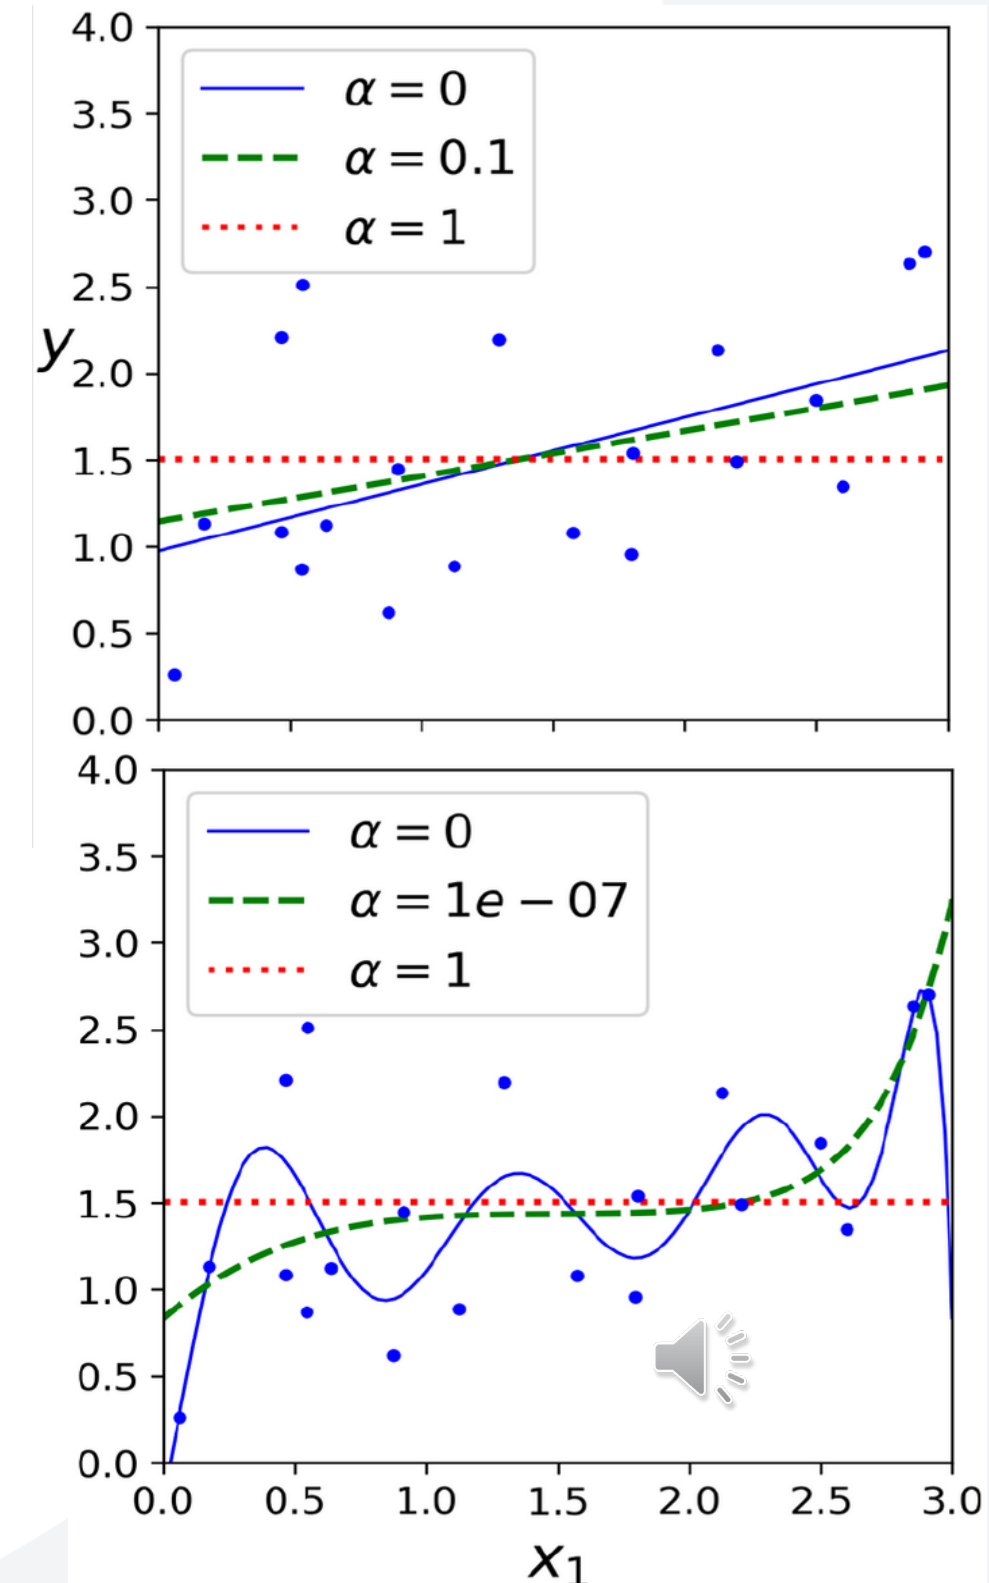
L1 (Lasso) Regularisation

- **LASSO** = Least **Absolute Shrinkage** and **Selection Operator**
- Effect: Penalises large model parameters to prevent overfitting. Tries to eliminate least important features ($\theta_i = 0$)
- Sparsity: Encourages sparsity in the model by setting some coefficients = 0, aiding in feature selection/interpretability.

$$J(\boldsymbol{\theta}) = \text{MSE}(\boldsymbol{\theta}) + \alpha \sum_{i=1}^n |\theta_i|$$

L1 Norm

- Scaling of Data: Important for setting α .
- Scikit-Learn: Use `penalty='l1'` parameter.
- Optimisation: Not differentiable at 0, but most optimisers can handle it.



Elastic Net

$r = 0$: L2 term dominates
 $r \sim 1$: L1 term dominates

A mixture of L1 and L2 regularisation.

$$J(\boldsymbol{\theta}) = \underbrace{\text{MSE}(\boldsymbol{\theta})}_{\text{Cost}} + \underbrace{r\alpha\sum_{i=1}^n |\theta_i|}_{\text{Data term}} + \underbrace{\frac{1-r}{2}\alpha\sum_{i=1}^n \theta_i^2}_{\text{Regularisation term}}$$

$\underbrace{\hspace{10em}}_{\text{Mixture ratio } r}$

$\underbrace{\hspace{10em}}_{\text{L1}} \quad \underbrace{\hspace{10em}}_{\text{L2}}$

Feature Selection: Can eliminate less important features.

Handling Collinearity: Often better than Lasso (L1) alone.

Penalty Term: Combination of absolute and squared coefficients.

Model Complexity: Balances between sparsity and retaining all features.

Which answer is correct?

What does L1 regularization (Lasso) encourage in a model?

- Inclusion of more features
- Large coefficients
- Sparsity by setting some coefficients to exactly zero
- Higher training error




Logistic Regression




Logistic Regression

- **Logistic regression:** a statistical method used to estimate the probability of an outcome (i.e. an event occurring or not occurring) based on one or more independent variables.
- Unlike linear regression, which outputs continuous values, logistic regression outputs a **probability** value **between 0 and 1**.

Linear Regression



Logistic Regression




Logistic Regression

- Combines a linear model with a logistic function (sigmoid)

$$\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \cdots + \theta_n x_n = \mathbf{x}^T \cdot \boldsymbol{\theta} \quad \text{Linear model}$$

$$\sigma(t) = \frac{1}{1 + \exp(-t)} \quad \text{Logistic function (sigmoid)}$$

$$\hat{p} = h_{\boldsymbol{\theta}}(\mathbf{x}) = \sigma(\mathbf{x}^T \boldsymbol{\theta}) \quad \text{Logistic model (regression)}$$



- Logistic function (sigmoid, $\sigma(t)$) maps the linear combination of input features ($\mathbf{x}^T \cdot \boldsymbol{\theta}$; expressed as log-odds) to a probability value (p) between 0 and 1.



Logistic Regression

- Combines a linear model with a logistic function (sigmoid)

$$\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \dots + \theta_n x_n = \mathbf{x}^T \cdot \boldsymbol{\theta} \quad \text{Linear model}$$

$$\sigma(t) = \frac{1}{1 + \exp(-t)} \quad \text{Logistic function (sigmoid)}$$

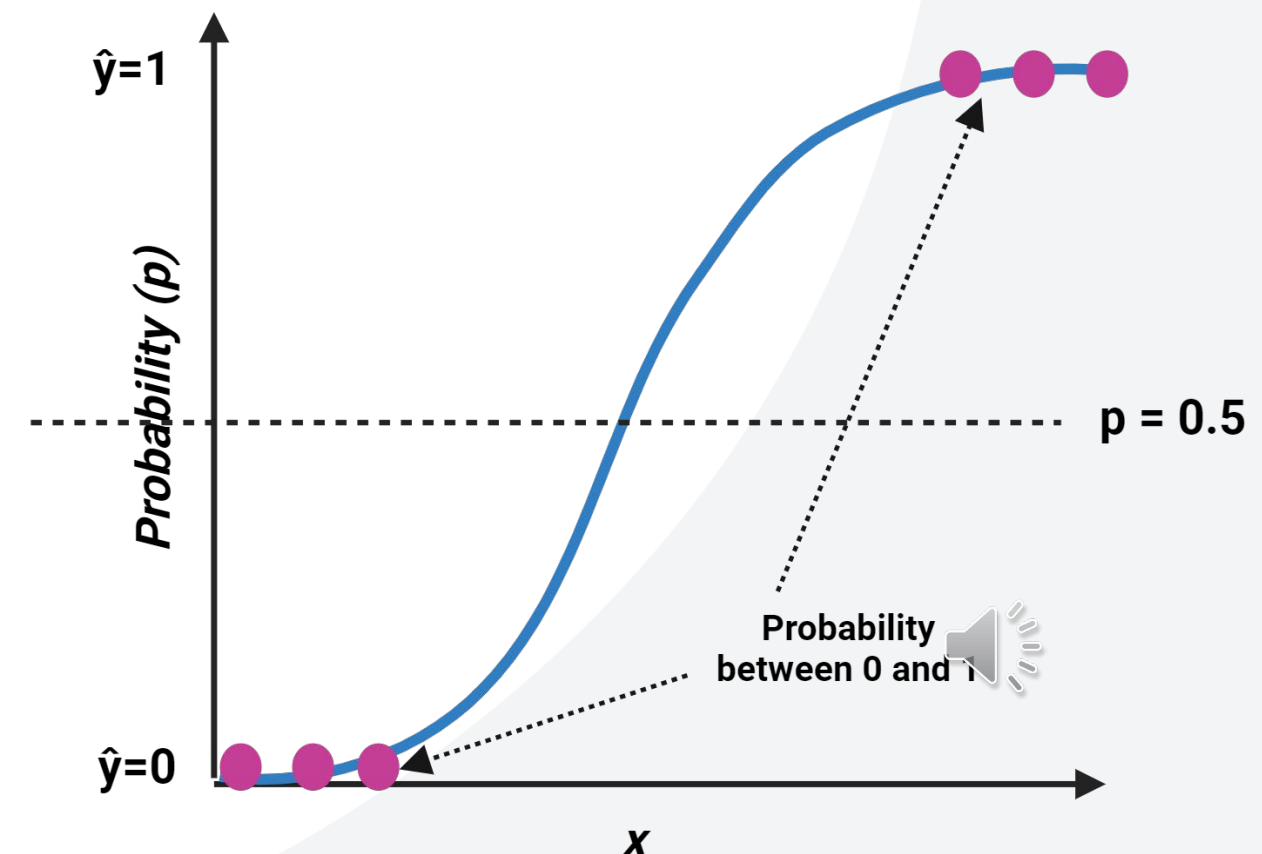
$$\hat{p} = h_{\boldsymbol{\theta}}(\mathbf{x}) = \sigma(\mathbf{x}^T \boldsymbol{\theta}) \quad \text{Logistic model (regression)}$$

- Used for binary classification ML problems where the output is a probability that the given input point belongs to a certain class:

$$\hat{y} = \begin{cases} 0 & \text{if } \hat{p} < 0.5 \\ 1 & \text{if } \hat{p} \geq 0.5 \end{cases}$$

Apply threshold = binary classifier

- Negative inputs = $\sigma(t) < 0.5$
- Positive inputs = $\sigma(t) > 0.5$
- Gradual probability change instead of a hard boundary.



Logit Regression = Logistic Regression

- **Logit regression:** Another name for logistic regression. The term "logit" refers to the log-odds function, which is the **inverse** of the logistic function.

$$\text{logit}(p) = \log\left(\frac{p}{1-p}\right) \quad \text{Log-odds function}$$

- Transforms probabilities into log-odds (converse of logistic regression).
- Same as logistic regression = process of using the logistic function (sigmoid) to predict binary outcomes.

Training a logistic model

- To train a logistic model, we need to have an appropriate cost/loss function to optimise.
- Can be achieved by considering probabilities:
 - Try to maximise probability of observing targets y given data x :

Probability of observing binary data y if real probability is p :

$$P(y) = \begin{cases} p & \text{if } y = 1 \\ 1 - p & \text{if } y = 0 \end{cases} \quad \text{or} \quad P(y) = p^y (1 - p)^{1-y}$$

For multiple events (set of observations), multiply the probabilities for each event:

$$P(\mathbf{y}) = \prod_i \hat{p}_i^{y_i} (1 - \hat{p}_i)^{1-y_i}$$

Take the log to simplify:

$$\log(P(\mathbf{y})) = \sum_i y_i \log(\hat{p}_i) + (1 - y_i) \log(1 - \hat{p}_i)$$



Training a logistic model

$$\log(P(\mathbf{y})) = \sum_i y_i \log(\hat{p}_i) + (1 - y_i) \log(1 - \hat{p}_i)$$


Our model predicts the estimated probability output by the regression

$$\hat{p} = h_{\boldsymbol{\theta}}(\mathbf{x}) = \sigma(\mathbf{x}^T \boldsymbol{\theta})$$

To train the model, we minimise the *negative log probability* (cost/loss function), also called **log loss** or **binary cross-entropy**:


$$J(\boldsymbol{\theta}) = -\frac{1}{m} \sum_{i=1}^m \left[y^{(i)} \log(\hat{p}^{(i)}) + (1 - y^{(i)}) \log(1 - \hat{p}^{(i)}) \right]$$

- Advantage: Cost function is convex, which is good for optimisation as it has no local minima



Logistic regression: Decision boundary

- **Decision boundary:** Helps decide which class a new data point belongs to, based on its features.
 - In logistic regression, the decision boundary is linear/monotonic, meaning it creates a straight line (in 2D) or a flat plane (in 3D) to separate classes.
 - Normal to the decision boundary is direction in which the probability of being in a specific class increases most (due to s-shaped sigmoid).




- **Near the boundary:** change in probability is more gradual.
- **Further from the boundary:** probability approaches 0 or 1 more quickly.



Logistic regression: Decision boundary

- **Probability prediction:** Graph below illustrates how the predicted probability changes with petal width.
- The intersection at $p=0.5$ is the decision boundary.
- This graph shows the rate of change near the decision boundary is the largest (probabilities plateau as you move further away).



Logistic regression with scikit-learn

Can have ways to predict both class and probabilities:

```
from sklearn.linear_model import LogisticRegression

log_model = LogisticRegression()
log_model.fit(X, y)
log_model.predict(X_new) # predict class
log_model.predict_proba(X_new) # predict probabilities
```

Softmax Regression

- Methodology so far explains the use of logistic regression for binary classification, but what if we have a multi-class classification problem?
- **Softmax:** Extends logistic regression framework to multiple classes. Also known as Multinomial Logistic Regression.


For each class k , we fit a linear model:

$$s_k(\mathbf{x}) = \mathbf{x}^\top \boldsymbol{\theta}^{(k)} \leftarrow$$

- $s_k(x)$: 'logits'
- \mathbf{x} : Input feature vector
- $\boldsymbol{\theta}^{(k)}$: Model parameters for class k

Softmax function: converts linear model outputs ('logits') for each class into probabilities:

$$\hat{p}_k = \sigma(\mathbf{s}(\mathbf{x}))_k = \frac{\exp(s_k(\mathbf{x}))}{\sum_{j=1}^K \exp(s_j(\mathbf{x}))}$$

- *exp* : Exponentiation makes all logits positive
- **Normalisation:** Sum of all probabilities $\sum_k \hat{p}_k = 1$ 
- **Range:** Each probability \hat{p}_k is between 0 and 1
- **Monotonic:** Larger logits lead to higher \hat{p}_k

Softmax Regression

Cost function = cross entropy, measures how well the predicted probabilities match the actual labels:


$$J(\Theta) = -\frac{1}{m} \sum_{i=1}^m \sum_{k=1}^K y_k^{(i)} \log(\hat{p}_k^{(i)})$$

Aim: Minimise this cost function ($J(\Theta)$).

- **Interpretation:** $y_k^{(i)}$ is 1 if the i -th sample belongs to class k , 0 otherwise.
- This cost function is differentiable, so we can use Gradient Descent as an optimisation scheme.

Comparison to logistic regression:

- Both fit a linear model to the input features and apply a non-linear transformation to produce probabilities (logistic regression - sigmoid function; Softmax – softmax function).
- Logistic Regression maps input features to a single probability value (binary classification); Softmax Regression maps input features to a probability distribution over multiple classes (multi-class classification).
- Both use some form of cross-entropy loss as cost function.




Support Vector Machines (SVM)




Support Vector Machines (SVMs)

- **Support Vector Machines (SVMs):** Type of supervised machine learning algorithm used for classification and regression tasks.
- Particularly effective for high-dimensional spaces and known for their ability to create clear boundaries between classes.
- **Key Idea:** Find the widest possible decision boundary between two classes while minimising errors (boundary violations).
- SVMs are based on distances between data points and the boundary.



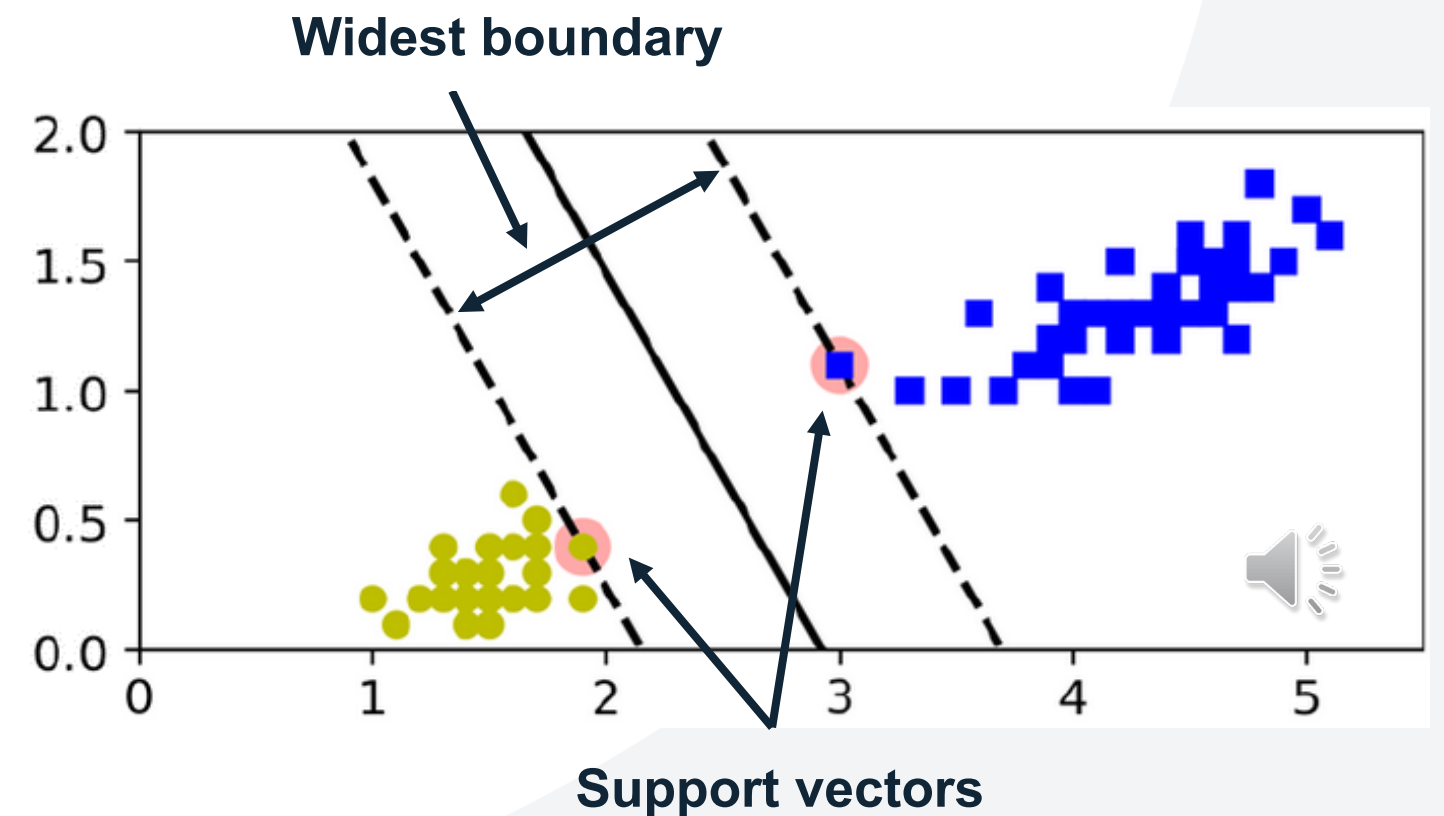
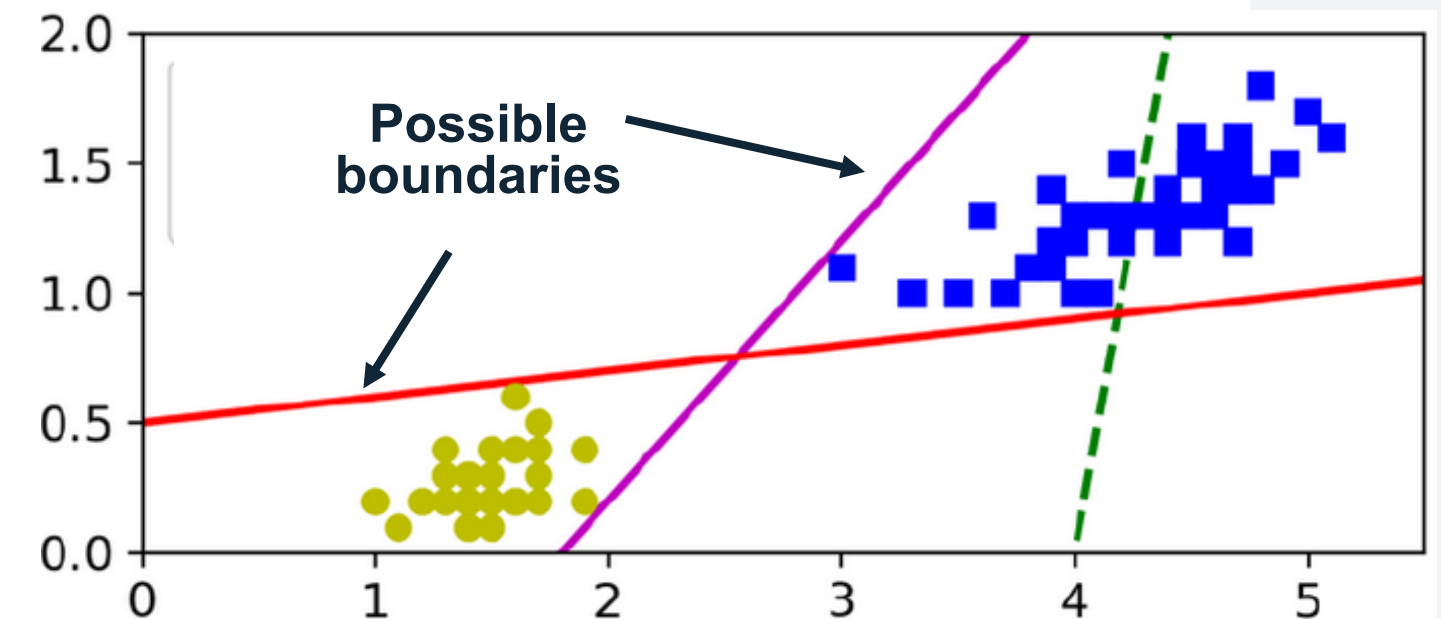
Support Vector Machines (SVM)

- **Support Vectors:** Boundary is defined by a few key data points called support vectors → Can identify boundary just from this small subset of samples.
- Good for complex boundaries, and small to medium datasets (large datasets can be slow).
- Boundary shape can be linear or non-linear.
- Uses a mathematical technique to called the 'kernel trick' handle non-linear boundaries efficiently.



Linear SVM – Hard Margin

- Decision boundary is straight line (2D) or linear plane/hyperplane (higher dimensions).
- **Zero Boundary Violations:** This means no data points from one class are allowed to fall on the other side of the boundary.
- **Sensitive to Scaling:** Because SVMs use distances, it's crucial that all input features are on a similar scale. If features have different scales, the SVM might not work correctly.




Linear SVM – Soft Margin

- **Allow Some Boundary Violations:** Allow some data points to be on the wrong side of the boundary to handle cases where classes are not perfectly separable.
 - Balance between having a wider boundary and allowing some violations (misclassified points).
 - The parameter C determines how much we allow boundary violations.
- **Small C (Near Zero):** Focusing on maximising the boundary width with minimal importance given to violations.
 - **Large C :** Achieves a narrower boundary and makes the model stricter about violations.


Wider boundary, more violations

$C = 1$



Less violations, narrower boundary

$C = 100$




Hinge loss function & SVM

- Can write the SVM loss using the Hinge loss function:

$$\max(0, 1 - t)$$

where t is a value that combines the model's prediction and the actual class label.



When $t \geq 1$:

Model's prediction is correct and confident.

Hinge Loss = 0: No penalty.

When $0 < t < 1$:

Model's prediction is correct but not confident enough.

Hinge Loss > 0: There's a small penalty. Smaller penalty the closer t is to 1.

When $t \leq 0$:

Model's prediction is wrong.

Hinge Loss > 0: There's a significant penalty. Larger the penalty the further t is from 0.



- You will see this as an option in other machine learning methods.

Non-Linear Boundaries

- **Problem:** Many decision boundaries are not linear
- **Idea:** Add non-linear functions of the input data as extra features and use a linear boundary in this higher dimensional space
 - Longer computation time

i.e. Gaussian Radial Basis Function (RBF)


Distance between samples

$$\phi_{\gamma}(\mathbf{x}, \ell) = \exp \left(-\gamma \overbrace{\|\mathbf{x} - \ell\|^2}^{\text{Distance between samples}} \right)$$

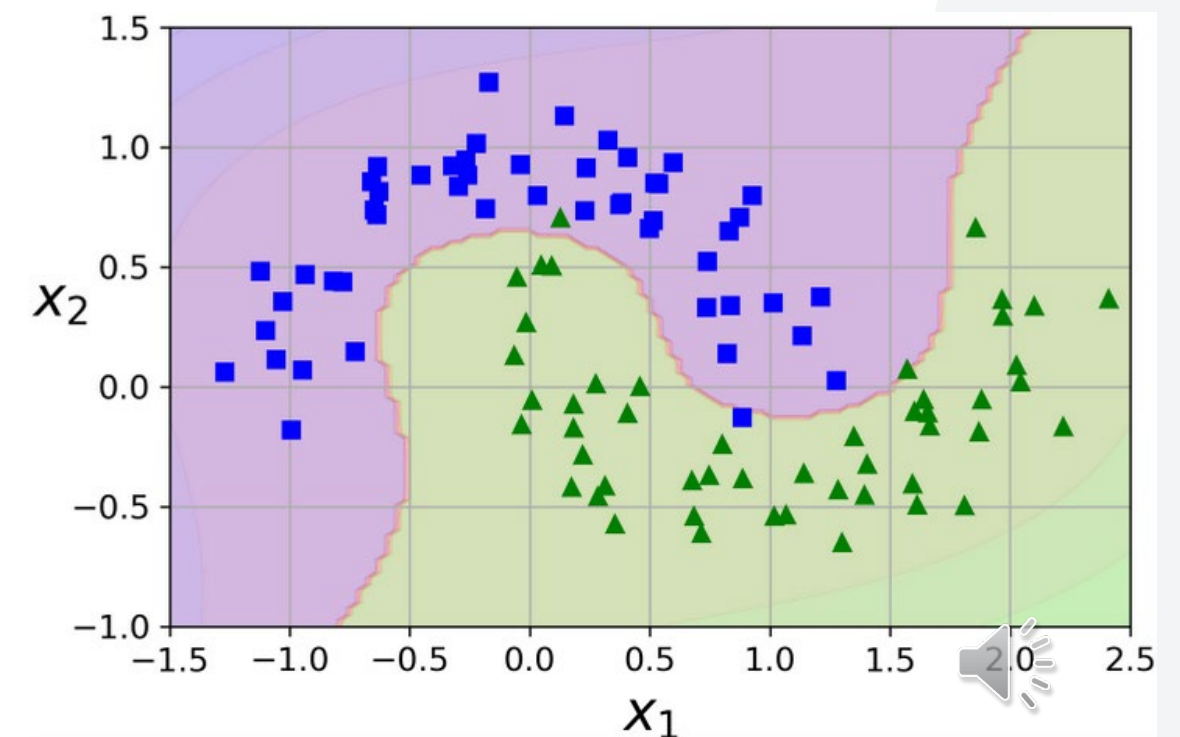

New non-linear feature

Hyperparameter Gamma

Not separable ☹️



Separable 😊

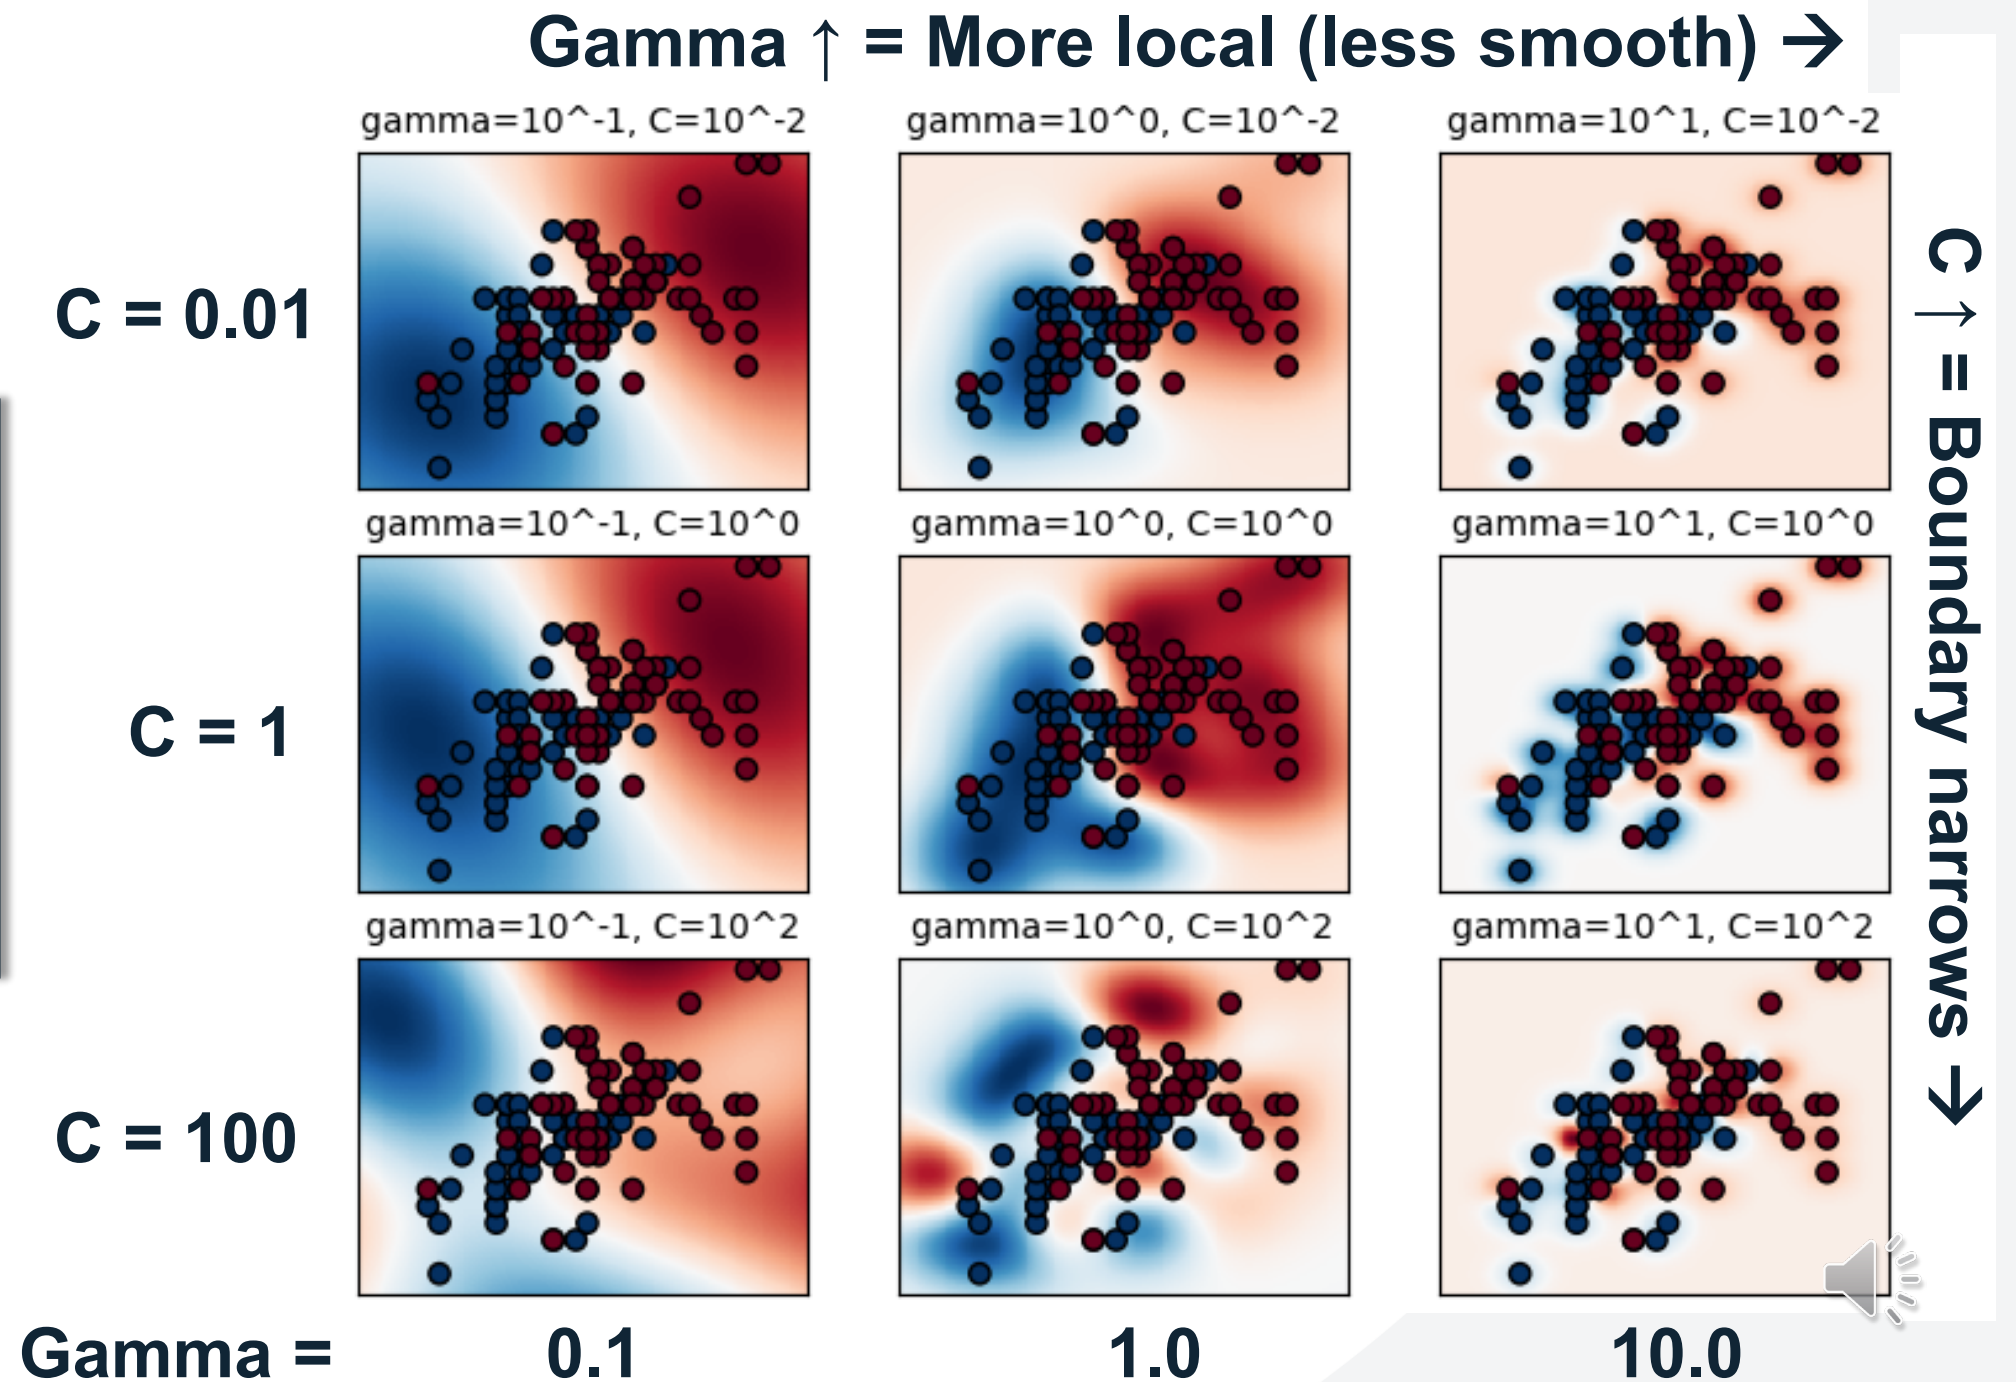


From Geron, Hands On ML

SVM Hyperparameters: C & Gamma

Example:

- 2 classes
- 50 samples each
- 2 (original) features
- Gaussian RBF kernel



Kernel 'trick'

- When we try to make our classifier more powerful, we sometimes add many extra features.
 - Can make model slow because of the large number of calculations needed.
- **Kernel trick:** Instead of calculating these extra features directly, we can use a clever shortcut - the SVM only needs to know the dot product between pairs of samples in the high-dimensional feature space, not the actual high-dimensional features themselves.



Kernel 'trick'

- **How it Works:**
 - Dot Products: Measure similarity between pairs of samples in a high-dimensional space.
 - Kernel Function: $K(a,b)$ specifies the dot product calculation.
 - Efficiency: Store dot products in a $N_{\text{sample}} \times N_{\text{sample}}$ matrix. Avoids need for explicit extra feature computation.

Common kernels

Linear: $K(a, b) = a^T b$

Linear classifier

Polynomial: $K(a, b) = (\gamma a^T b + r)^d$

Use polynomial feature combinations without having to compute them

Gaussian RBF: $K(a, b) = \exp(-\gamma ||a - b||^2)$

Use similarity to support vectors with Gaussian drop off

Sigmoid: $K(a, b) = \tanh(\gamma a^T b + r)$

Summary

- **Regularisation continued**
- **Logistic regression**
 - Probability-based method
 - Uses sigmoid (logistic) function to get outputs in $[0,1]$ range
 - Sigmoid & Softmax functions used in many deep learning models
- **Support Vector Machines**
 - Fit the widest possible boundary between two classes while limiting boundary violations
 - Create extra non-linear features \rightarrow high dimensional space where they can be separated with a linear boundary in this space
 - Best to try simple kernels first, compare different kernels
 - Hyperparameters also need to be optimised
 - Data needs to be scaled uniformly (internally it is distance-based)

Questions?

dhani.dharmaprani@adelaide.edu.au

