

# Using Machine Learning Tools PG

Week 9 – Convolutional Neural  
Networks

COMP SCI 7317

Trimester 2, 2024



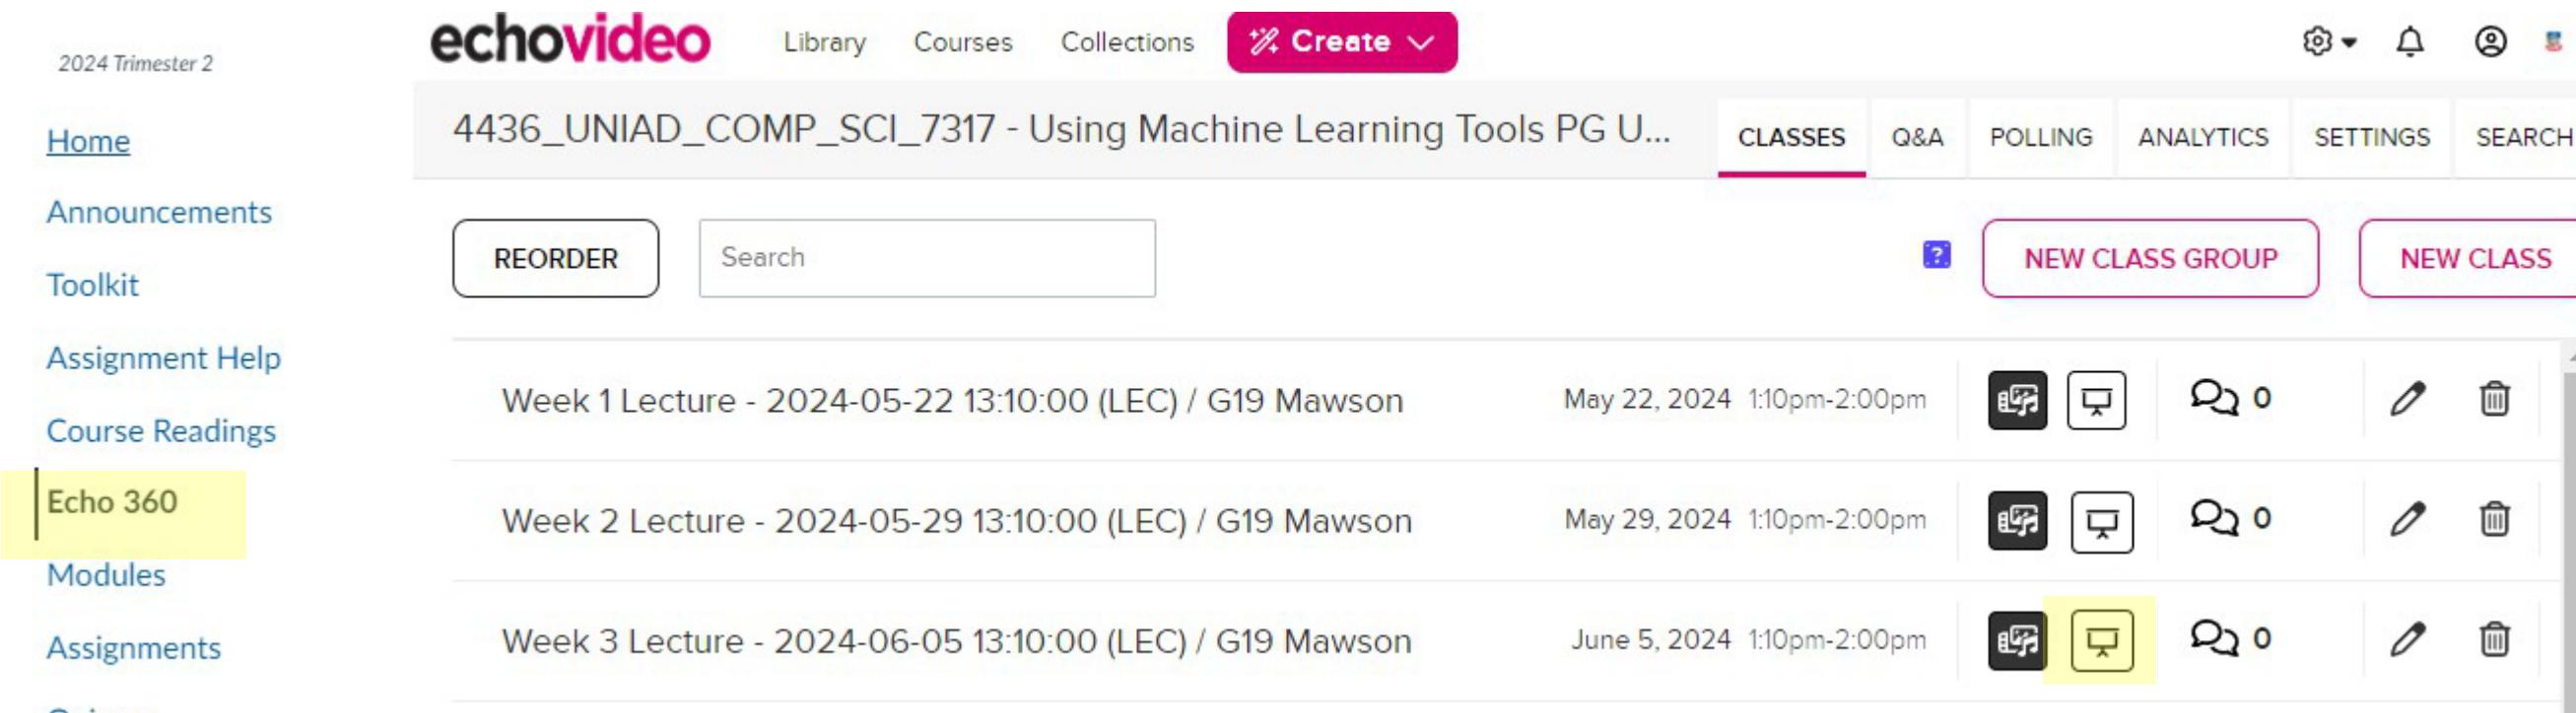
THE UNIVERSITY  
*of*ADELAIDE

150  
YEARS

# Before we start...

There will be some **interactive elements** throughout this lecture.

Please participate by heading to **MyUni > 4436\_COMP\_SCI\_7317 > Echo360 > Lecture Week 9** > 



The screenshot shows the EchoVideo interface for the course 4436\_UNIAD\_COMP\_SCI\_7317. The left sidebar includes links for Home, Announcements, Toolkit, Assignment Help, Course Readings, Echo 360 (highlighted in yellow), Modules, Assignments, and Quizzes. The main area displays three lectures:

Lecture	Date	Time	Location	Actions
Week 1 Lecture - 2024-05-22 13:10:00 (LEC)	May 22, 2024	1:10pm-2:00pm	G19 Mawson	0
Week 2 Lecture - 2024-05-29 13:10:00 (LEC)	May 29, 2024	1:10pm-2:00pm	G19 Mawson	0
Week 3 Lecture - 2024-06-05 13:10:00 (LEC)	June 5, 2024	1:10pm-2:00pm	G19 Mawson	0



THE UNIVERSITY  
of ADELAIDE

15C  
YEARS

# Last week... Neural Networks & Deep Learning

## 1. Components of a Neural Network

- Neurons: non-linearities and activation functions
- Network architecture
- Parameters: connections, weights and biases
- Number of parameters
- Deep Neural Networks

## 2. Loss functions, epochs, batches, optimisers and training

## 3. Classification and regression variants



THE UNIVERSITY  
of ADELAIDE

15C  
YEARS

# This week

- 1. Convolutional Neural Networks**
- 2. CNN Structures**
- 3. Pooling + Parameters**
- 4. Applications + Considerations**



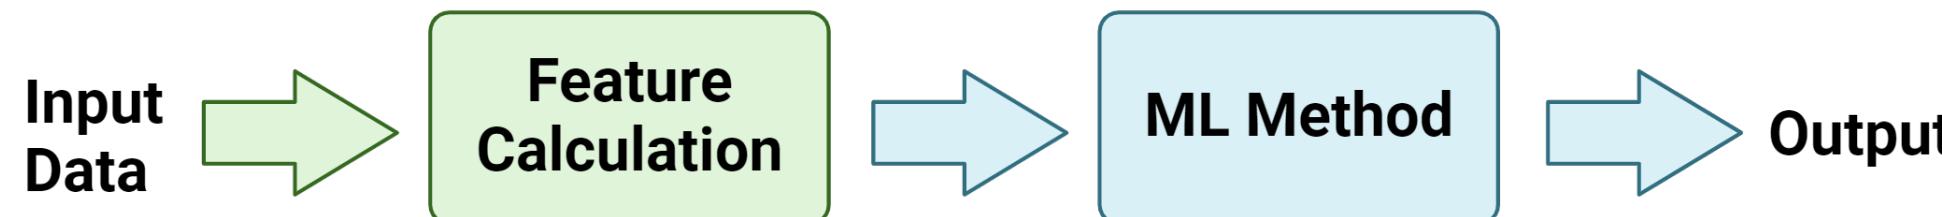
THE UNIVERSITY  
of ADELAIDE

**15C**  
YEARS

# Convolutional Neural Networks

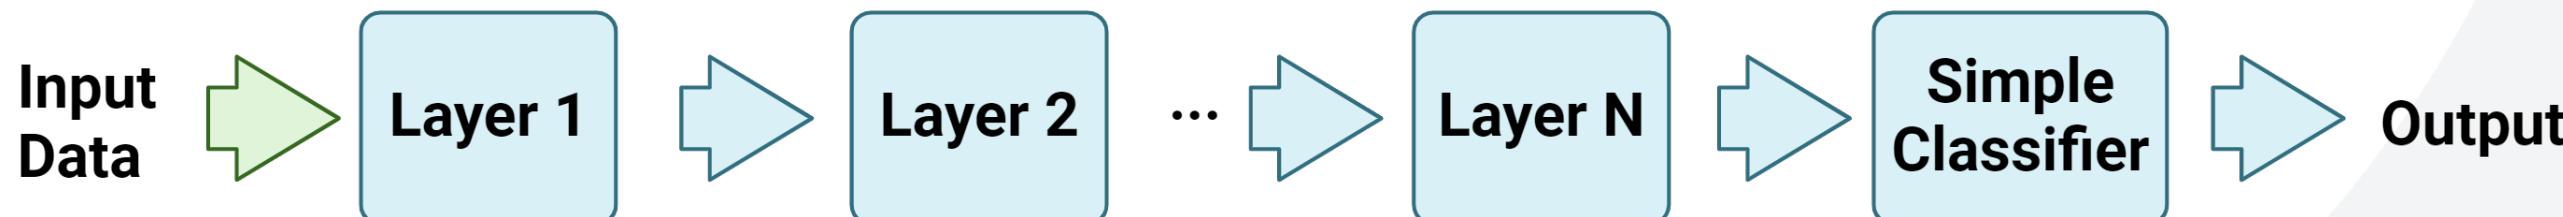
# Introduction to Neural Networks

## Traditional Machine Learning



- “Hand crafted features” i.e. pre-define the relevant features

## Deep Learning



- “Learned features”: Network automatically learns relevant features
- Tend to have many more parameters + need much more data



THE UNIVERSITY  
of ADELAIDE

150  
YEARS

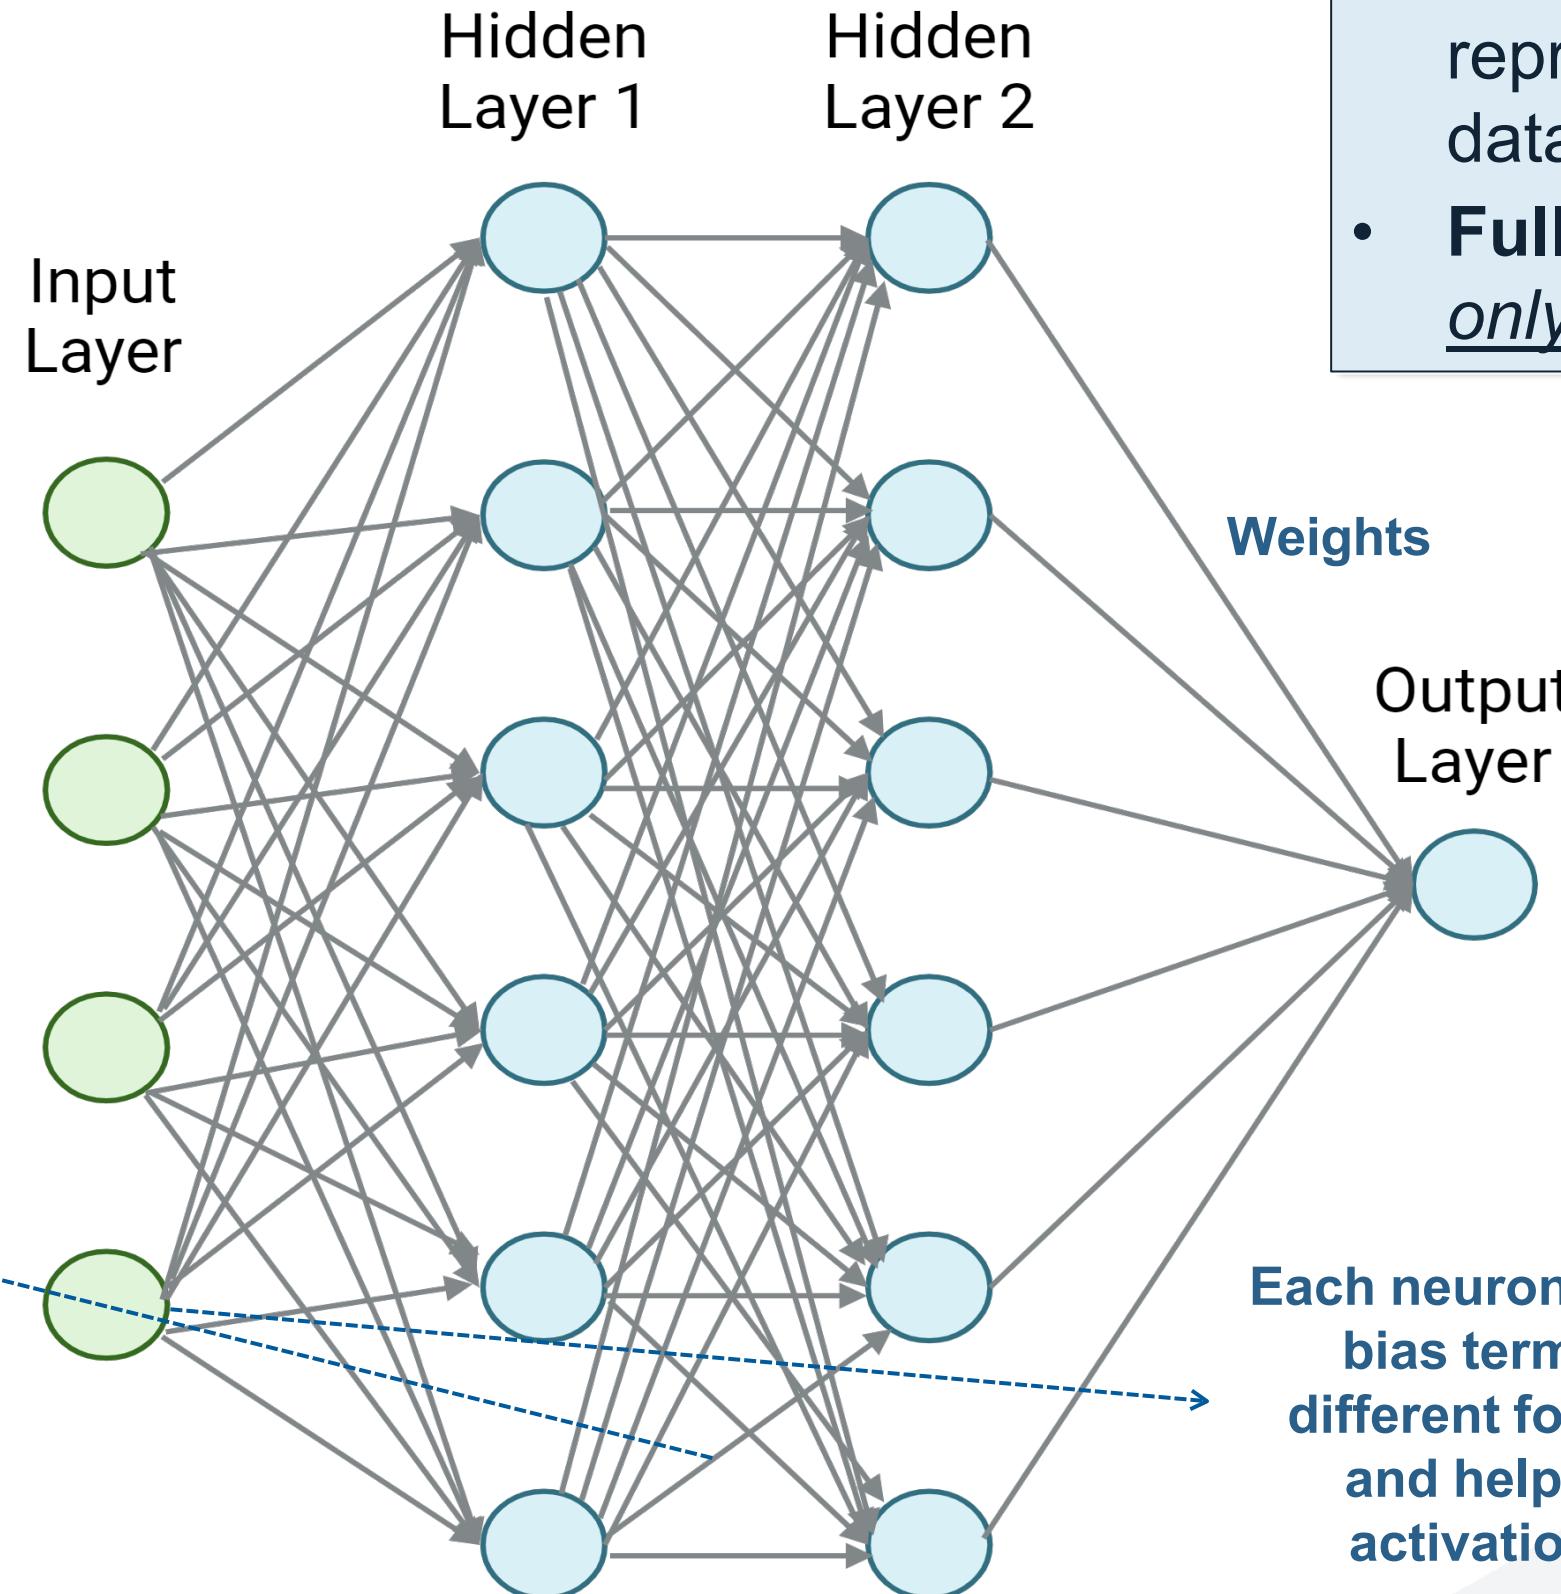
# Neural Network

## Input values as 1D vector

# Fully Connected Layout

$\delta(f)$

# Activation function



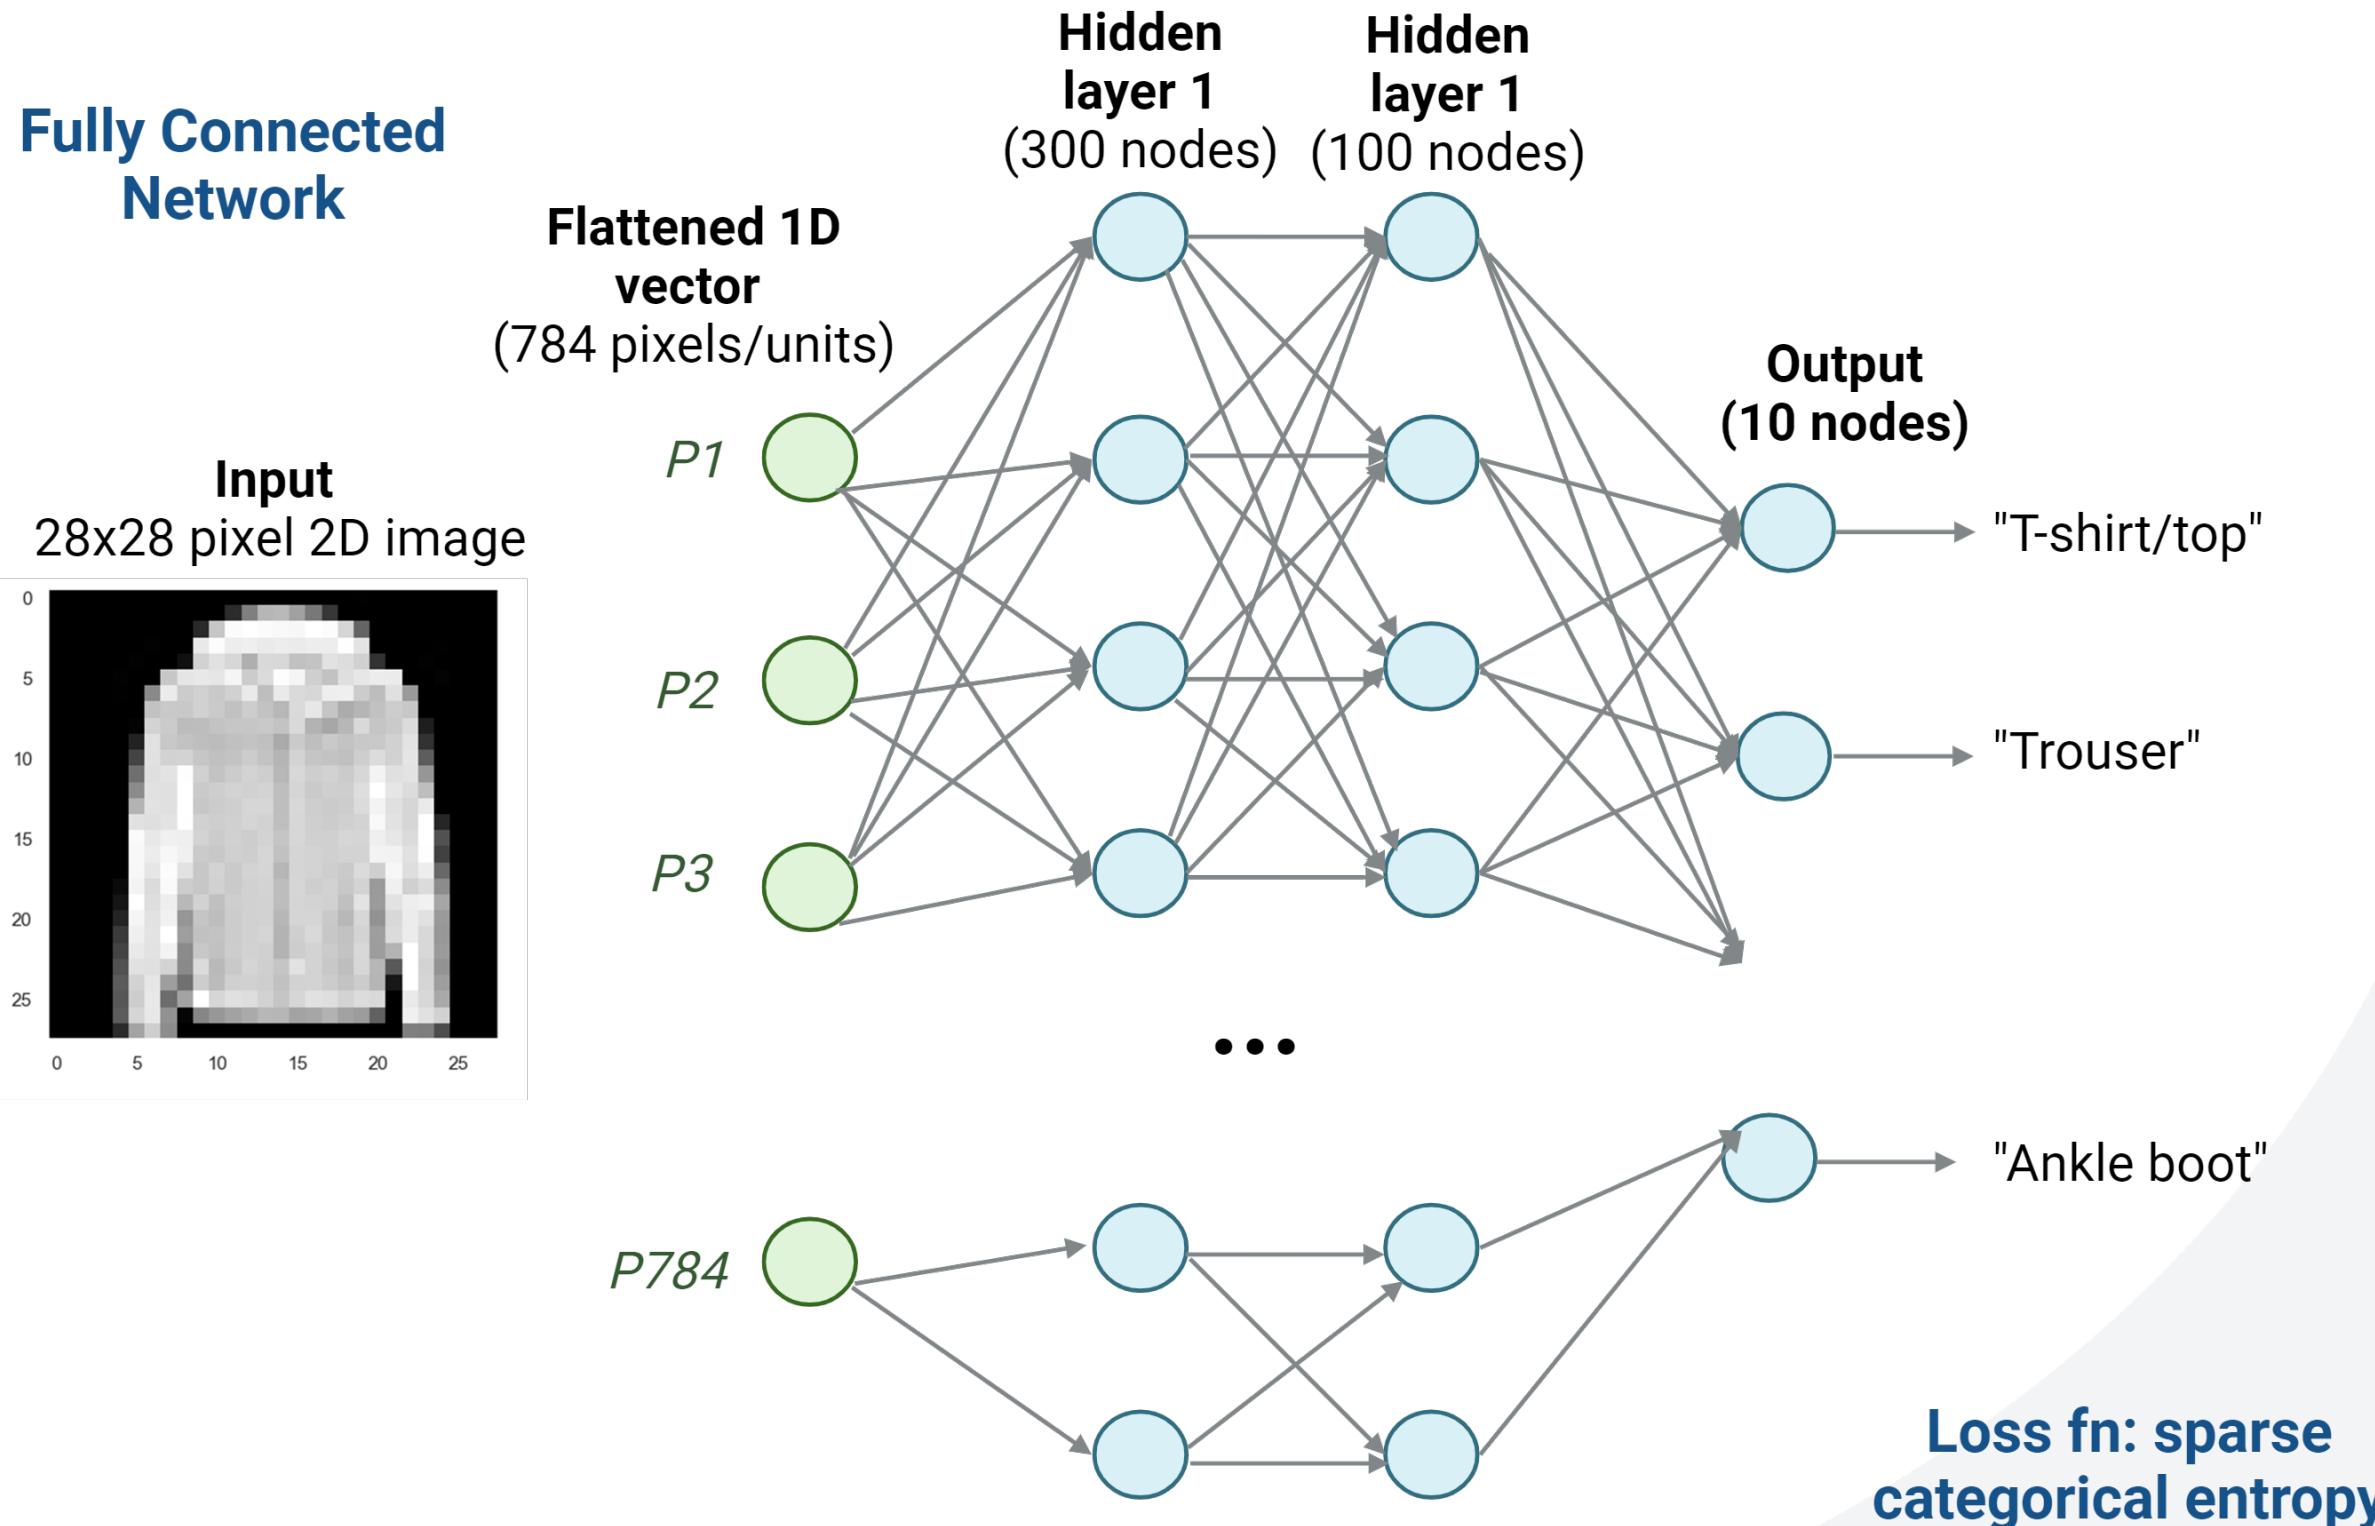
- Commonly used architecture.
  - Each green circle/input node represents a single value (i.e. data sample, pixel).
  - **Fully connected:** Each layer only connects to next layer.

Each neuron has an implicit bias term  $b$ , which is different for each neuron and helps adjust the activation threshold



**150** YEARS

# Neural Network: Image ‘Flattening’/Reshaping

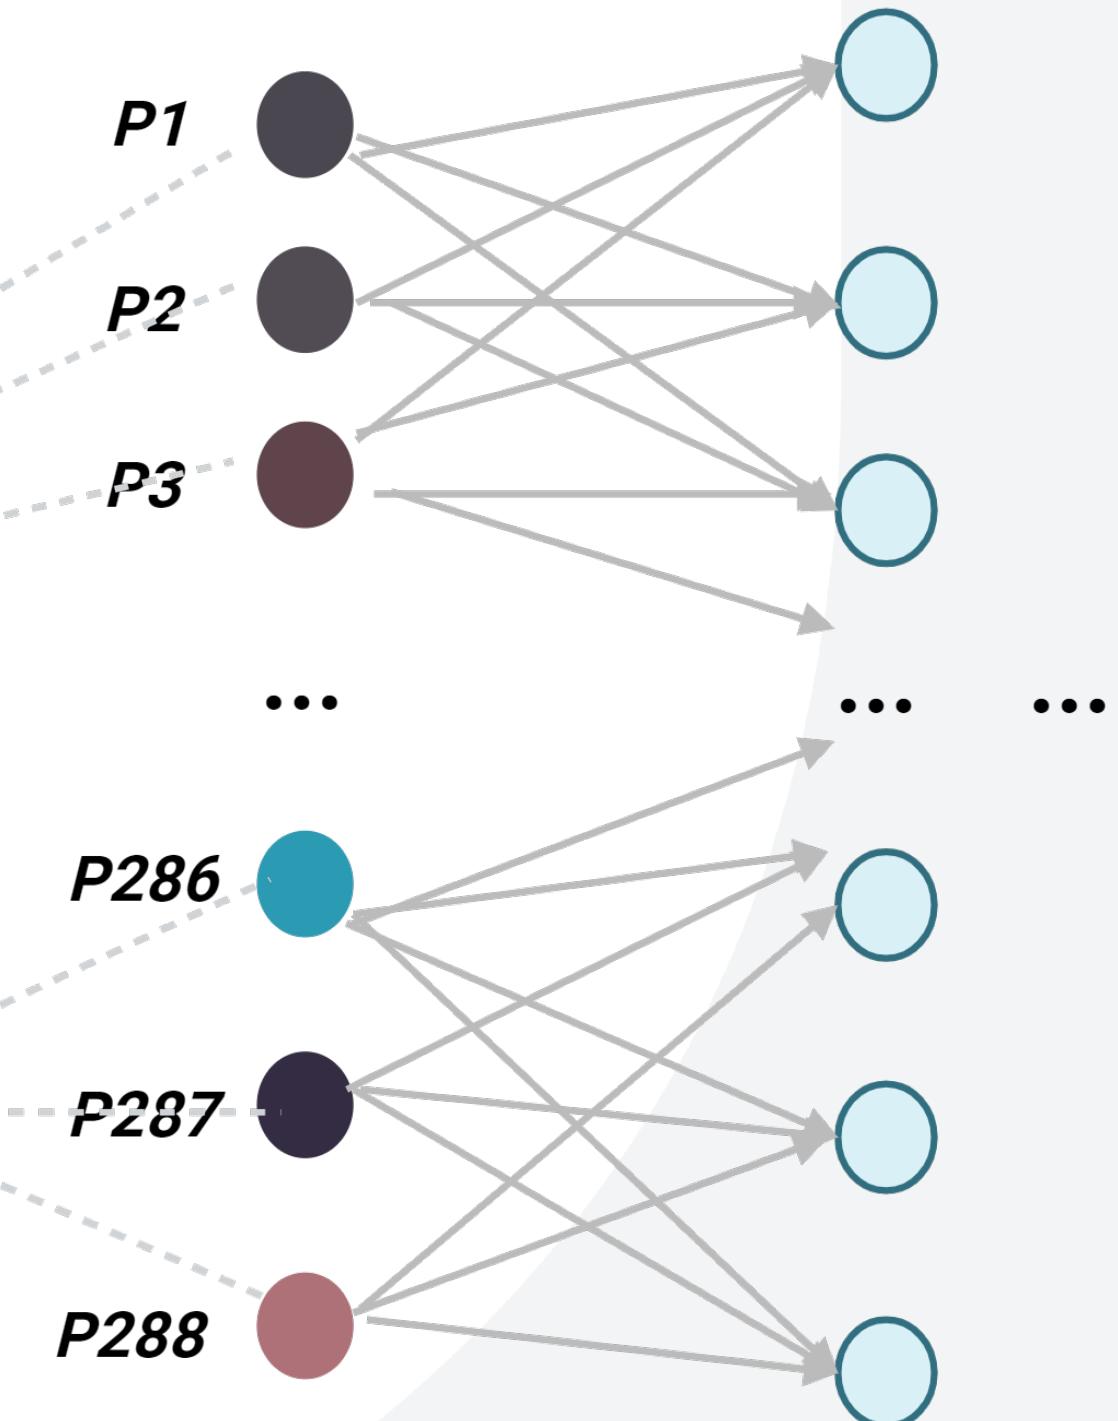
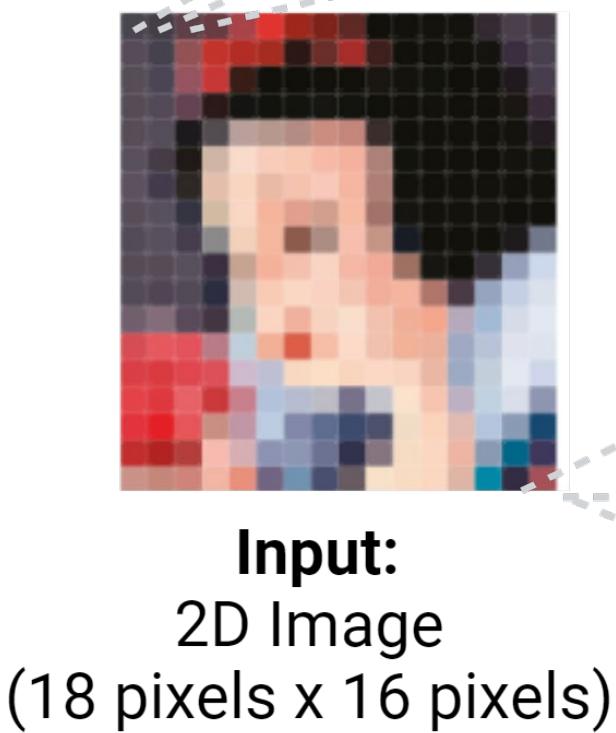


# Loss of spatial information

Network loses spatial information about the image. Relationships between adjacent pixels, which could represent important features like edges or textures, are not directly encoded in the input.

**Impact:** Network may struggle to learn spatial patterns, as it treats each pixel independently rather than considering their spatial relationships.

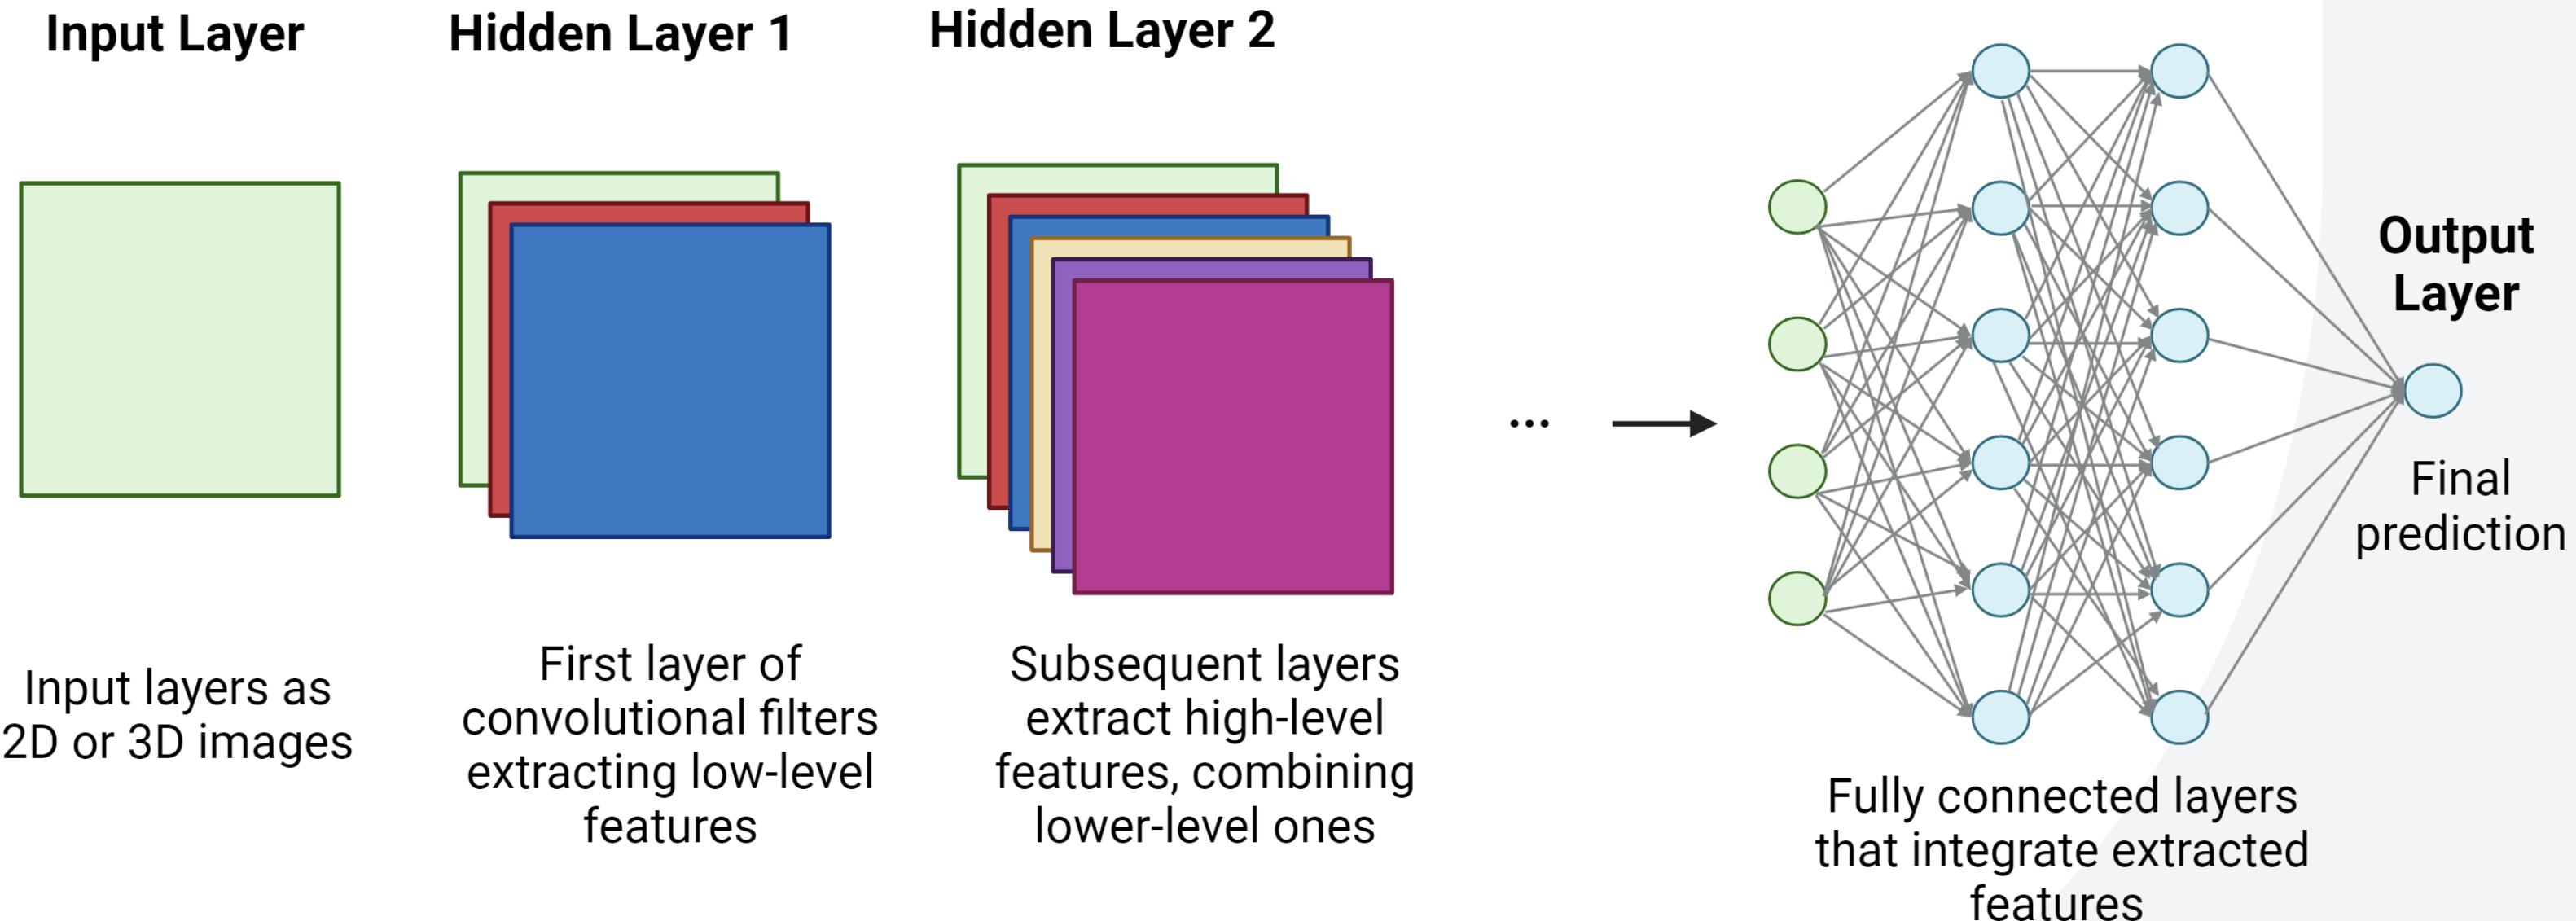
**Solution:** Use a linear operation that works with images and few parameters → **Convolution**



# Convolutional Neural Network

- Implicit bias term  $b$  not shown
- Activation function still applied

***Each layer only connects to next layer***

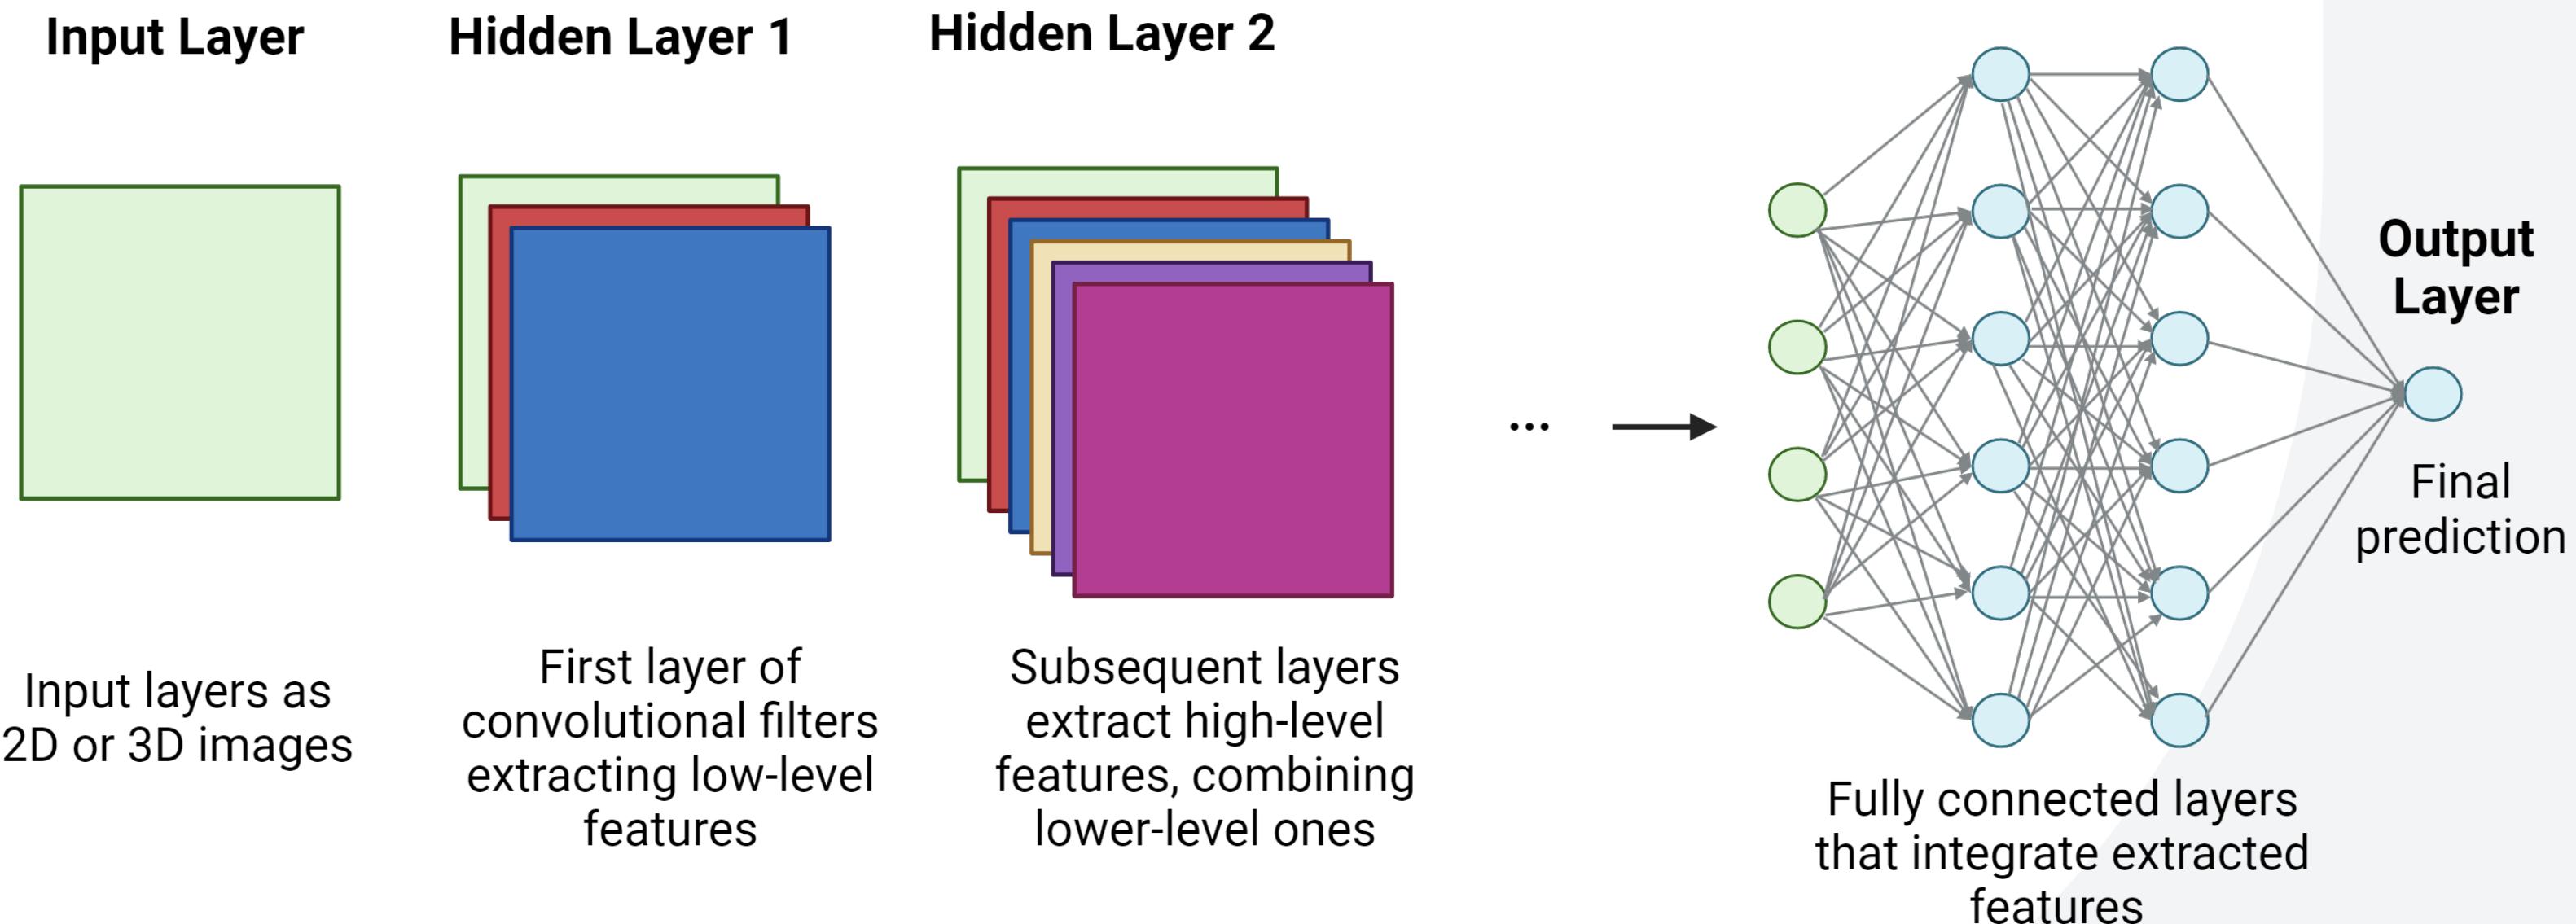


- A type of deep learning algorithm designed to process structured grid data such as images.
- Primarily used for image recognition, classification, and processing.

# Convolutional Neural Network

- Implicit bias term  $b$  not shown
- Activation function still applied

***Each layer only connects to next layer***



- Convolutional and pooling layers act as feature extractors.
- Fully connected layers act as a classifier based on the extracted features.

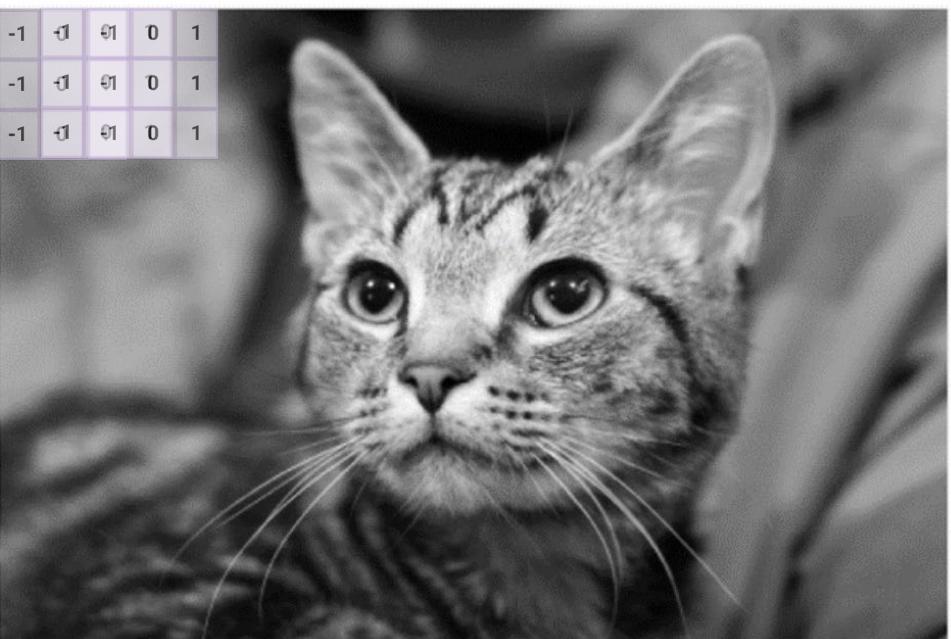
# CNN Structures

# Convolution

**Convolution:** A process that uses a small matrix (filter or kernel) to scan over an image and extract features. It involves sliding the filter over the input image, multiplying the overlapping values, and summing them to produce a new value in the output image.

- **Example:** edge detection (shown).

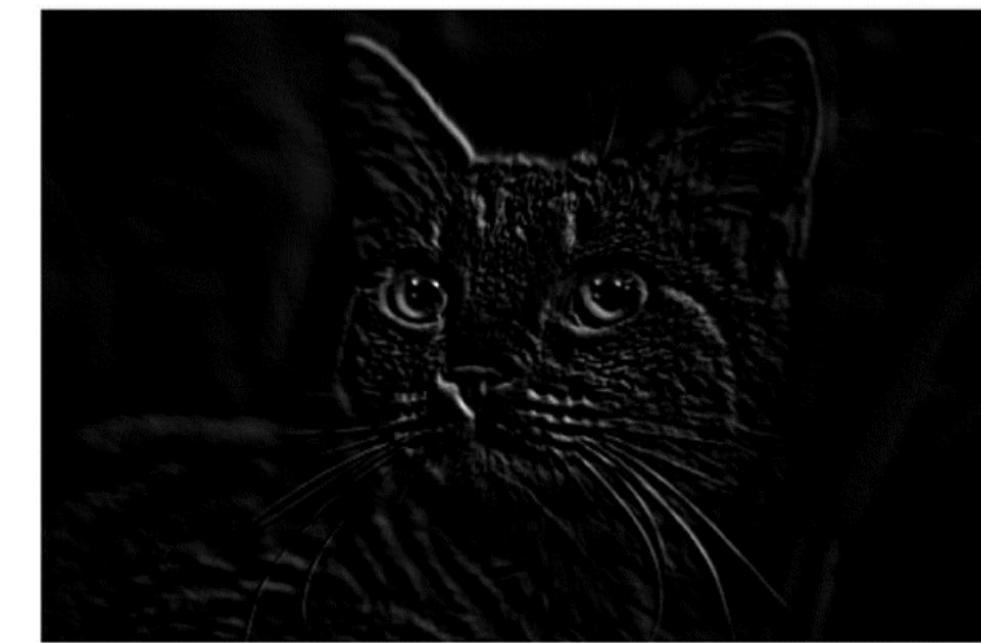
Original



Filter

-1	0	1
-1	0	1
-1	0	1

Convolved



Many different types of filters are available!

# Convolutional Layer

Image (5x5 pixels)

4	8	3	5	0
2	4	5	3	0
1	9	1	7	0
7	6	0	0	0
0	0	0	0	0

Zero values clustered together  
→ Spatial relationship

Filter

-1	2	-1
-2	4	-2
-1	2	-1

Generally chosen by NN

Output

0	19	-11	15	-13
-7	27	-15	22	-18
-6	38	-31	27	-17
9	26	-26	13	-7
8	5	-6	0	0

Convolutional layer uses a set of filters (also known as kernels) that slide over the input data. Output of each convolution produces a feature map

# Convolutional Layer

Image (5x5 pixels)

-1	2	-1			
-2	4	8	3	5	0
-1	2	4	5	3	0
1	9	1	7	0	
7	6	0	0	0	
0	0	0	0	0	

Filter

-1	2	-1
-2	4	-2
-1	2	-1

Output

0	19	-11	15	-13
-7	27	-15	22	-18
-6	38	-31	27	-17
9	26	-26	13	-7
8	5	-6	0	0

$$(4*4) + (8*-2) + (2*2) + (4*-1) = 0$$

Starting with centre value of the filter/kernel aligned with the top-left most pixel in the image

# Convolutional Layer

Image (5x5 pixels)

-1	2	-1			
-2	4	8	3	5	0
-1	2	4	5	3	0
1	9	1	7	0	
7	6	0	0	0	
0	0	0	0	0	

Filter

-1	2	-1
-2	4	-2
-1	2	-1

Values outside of the filter/kernel  
ignored → Same effect as 0 padding

Output

0	19	-11	15	-13
-7	27	-15	22	-18
-6	38	-31	27	-17
9	26	-26	13	-7
8	5	-6	0	0

$$(4*4) + (8*-2) + (2*2) + (4*-1) = 0$$

# Convolutional Layer

Image (5x5 pixels)

-1	2	-1			
-2	4	8	3	5	0
-1	2	4	5	3	0
1	9	1	7	0	
7	6	0	0	0	
0	0	0	0	0	

Filter

-1	2	-1
-2	4	-2
-1	2	-1

Alternative is to constrain filter/kernel to inside the image, but this makes output smaller as you can't align the centre with the edge pixel.

$$(4*4) + (8*-2) + (2*2) + (4*-1) = 0$$

Output

0	19	-11	15	-13
-7	27	-15	22	-18
-6	38	-31	27	-17
9	26	-26	13	-7
8	5	-6	0	0

# Convolution Layer

Image (5x5 pixels)

-1	2	-1		
4	8	3	5	0
2	4	5	3	0
1	9	1	7	0
7	6	0	0	0
0	0	0	0	0

Filter

-1	2	-1
-2	4	-2
-1	2	-1

Output

0	19	-11	15	-13
-7	27	-15	22	-18
-6	38	-31	27	-17
9	26	-26	13	-7
8	5	-6	0	0

$$(4 * -2) + (8 * 4) + (3 * -2) + (2 * -1) + (4 * 2) + (5 * -1) = 19$$

Moving filter/kernel so centre is aligned with next pixel in the image (to the right)

# Convolutional Layer

Image (5x5 pixels)

	-1	2	-1	
4	8	3	5	0
2	4	5	3	0
1	9	1	7	0
7	6	0	0	0
0	0	0	0	0

Filter

-1	2	-1
-2	4	-2
-1	2	-1

Output

0	19	-11	15	-13
-7	27	-15	22	-18
-6	38	-31	27	-17
9	26	-26	13	-7
8	5	-6	0	0

Repeat until the image is fully convolved with the filter/kernel

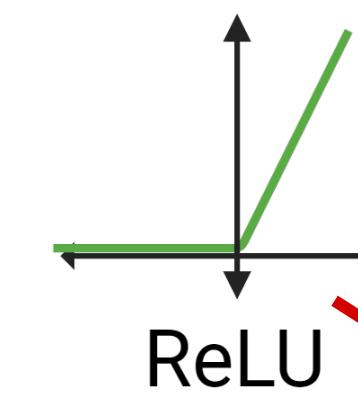
# Convolutional Layer

Image (5x5 pixels)

4	8	3	5	0
2	4	5	3	0
1	9	1	7	0
7	6	0	0	0
0	0	0	0	0

Filter

-1	2	-1
-2	4	-2
-1	2	-1



Output

0	19	-11	15	-13
-7	27	-15	22	-18
-6	38	-31	27	-17
9	26	-26	13	-7
8	5	-6	0	0

Applying convolution is a linear operation. To introduce non-linearity into the model to learn more complex patterns, an activation function (e.g., ReLU) is applied element-wise after convolution is performed.

i.e. **Convolution → Output feature map → Activation function**

# Convolutional Layer

Image (5x5 pixels)

4	8	3	5	0
2	4	5	3	0
1	9	1	7	0
7	6	0	0	0
0	0	0	0	0

Filter

-1	2	-1
-2	4	-2
-1	2	-1

Output

0+1	19+1	-11+1	15+1	-13+1
-7+1	27+1	-15+1	22+1	-18+1
-6+1	38+1	-31+1	27+1	-17+1
9+1	26+1	-26+1	13+1	-7+1
8+1	5+1	-6+1	0+1	0+1

*Bias*  
+1

We can also have bias terms, which helps the model learn an additional offset to improve its ability to fit the data (similarly to other neural networks). This adds a value to each element in output feature map.  
Example on previous slides equivalent to bias of 0.

# Which answer is correct?

What is the main purpose of an activation function in a neural network?

- To initialise weights
- To introduce non-linearity
- To normalise input data
- To optimise the loss function

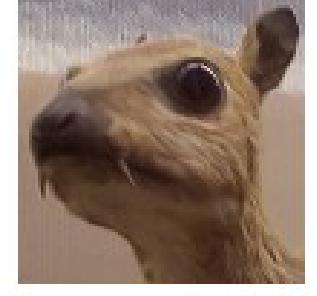
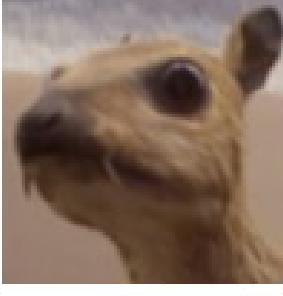


THE UNIVERSITY  
of ADELAIDE

15C  
YEARS

# Convolutional Layer

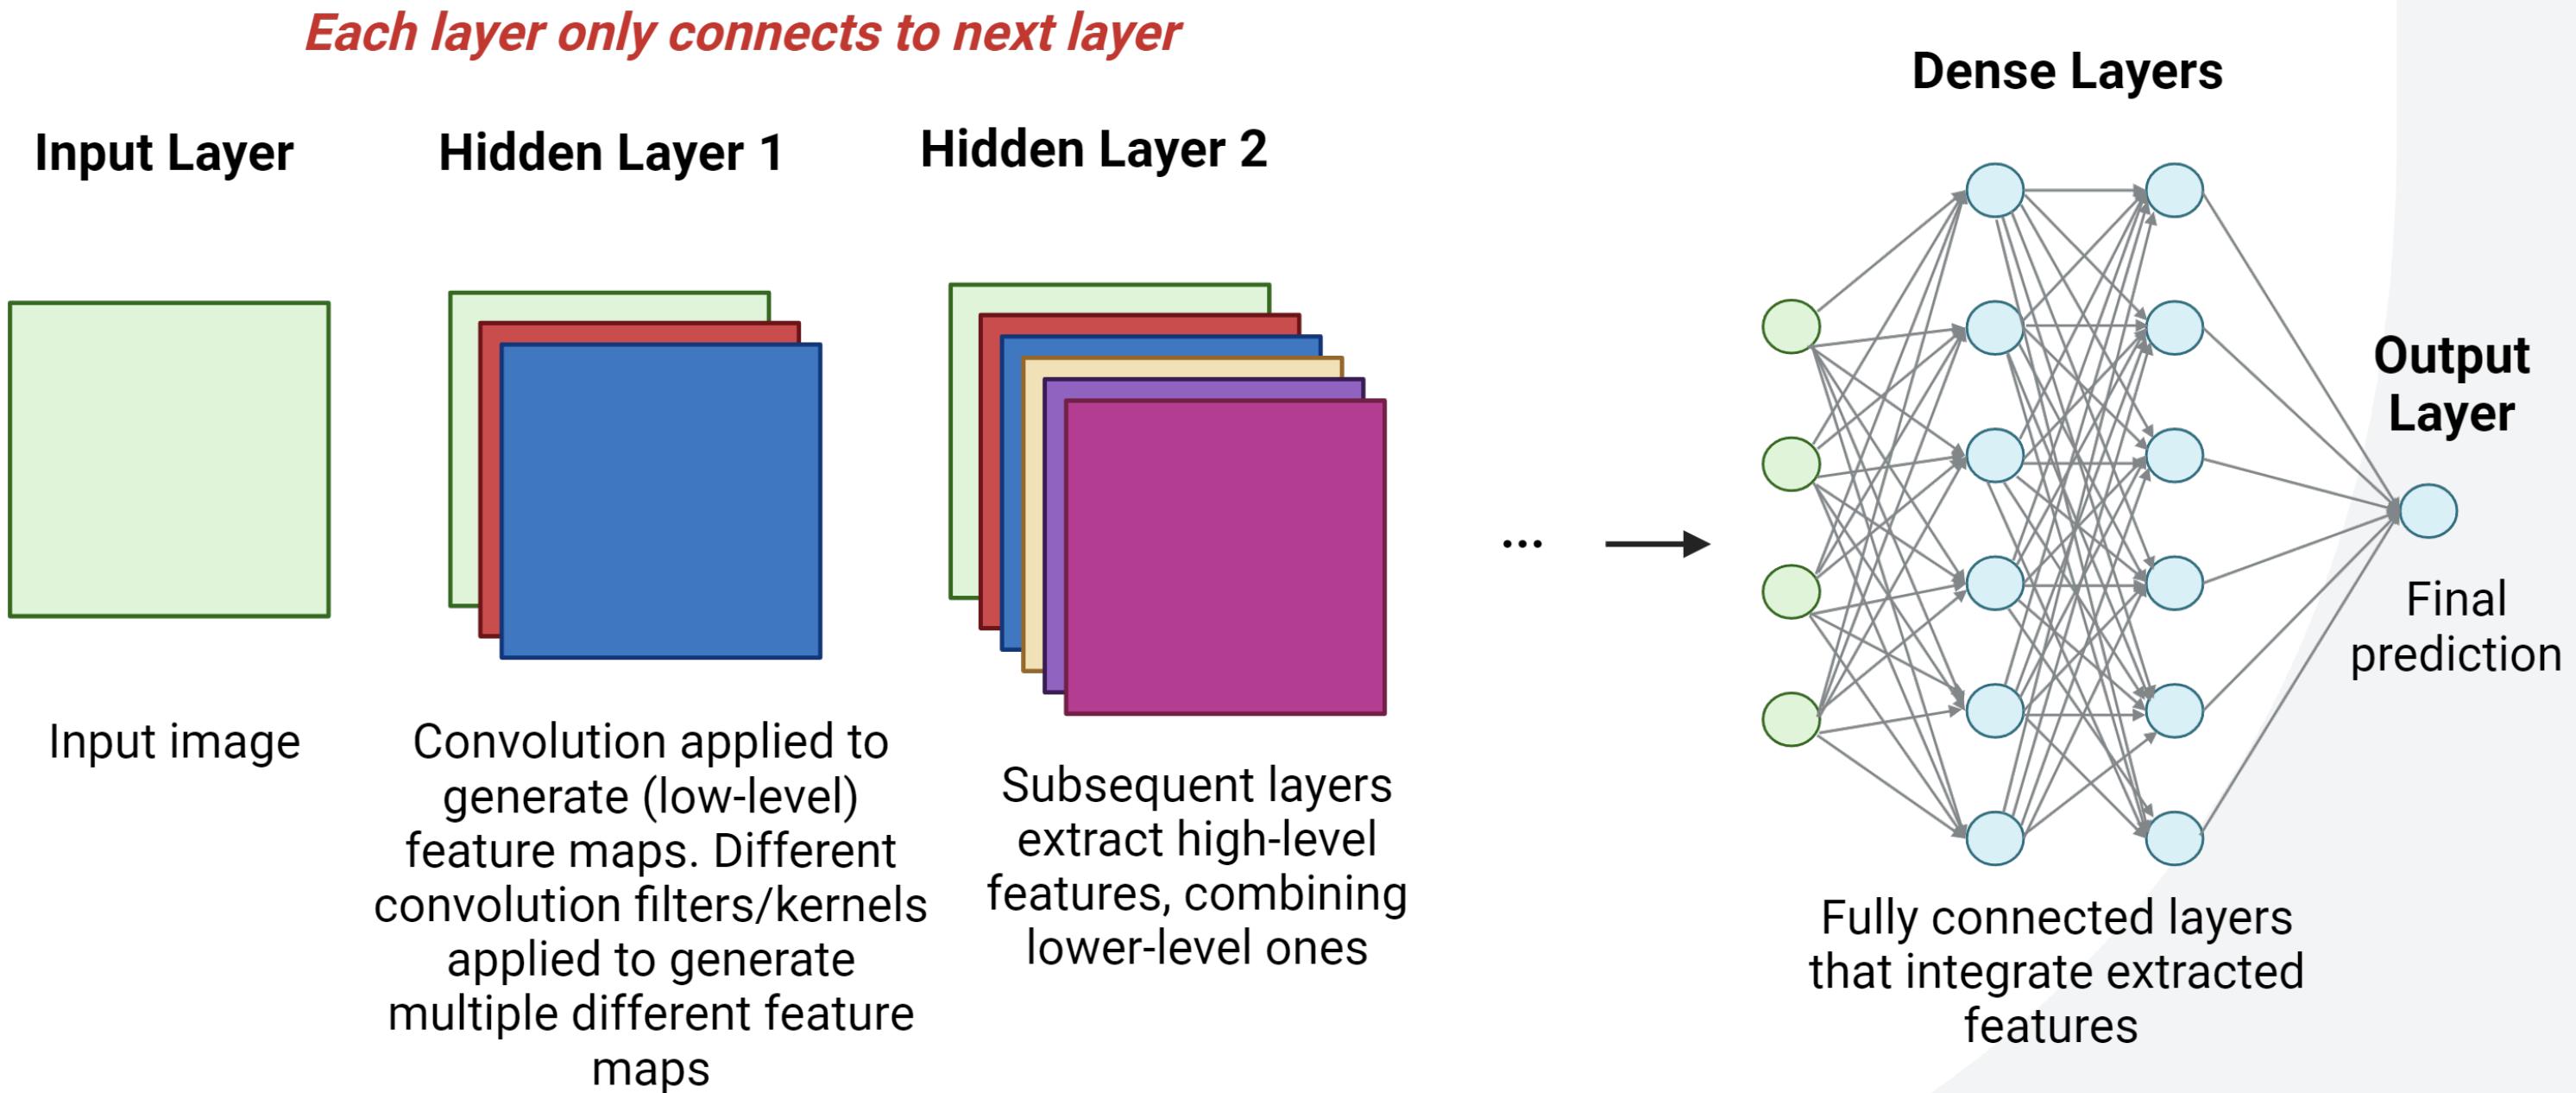
- Only uses a small portion of the image at a time, so tends to be quite computationally efficient operation.
- Many types of filters/kernels exist, depending on what is most suitable to extract features.
  - **Examples:** edge detectors, blur filters, etc.
- **Stride** determines how much the filter moves at each step (usually 1 for convolution), and **padding** helps in controlling the output size by adding extra pixels around the border of the input.
- Size of filter/kernel affects the granularity of the feature extraction (usually 3x3)
  - Smaller filters capture fine details
  - Larger filters capture broader patterns
  - Better to have smaller filters + more layers

Original	Gaussian Blur
$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$	$\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$
	

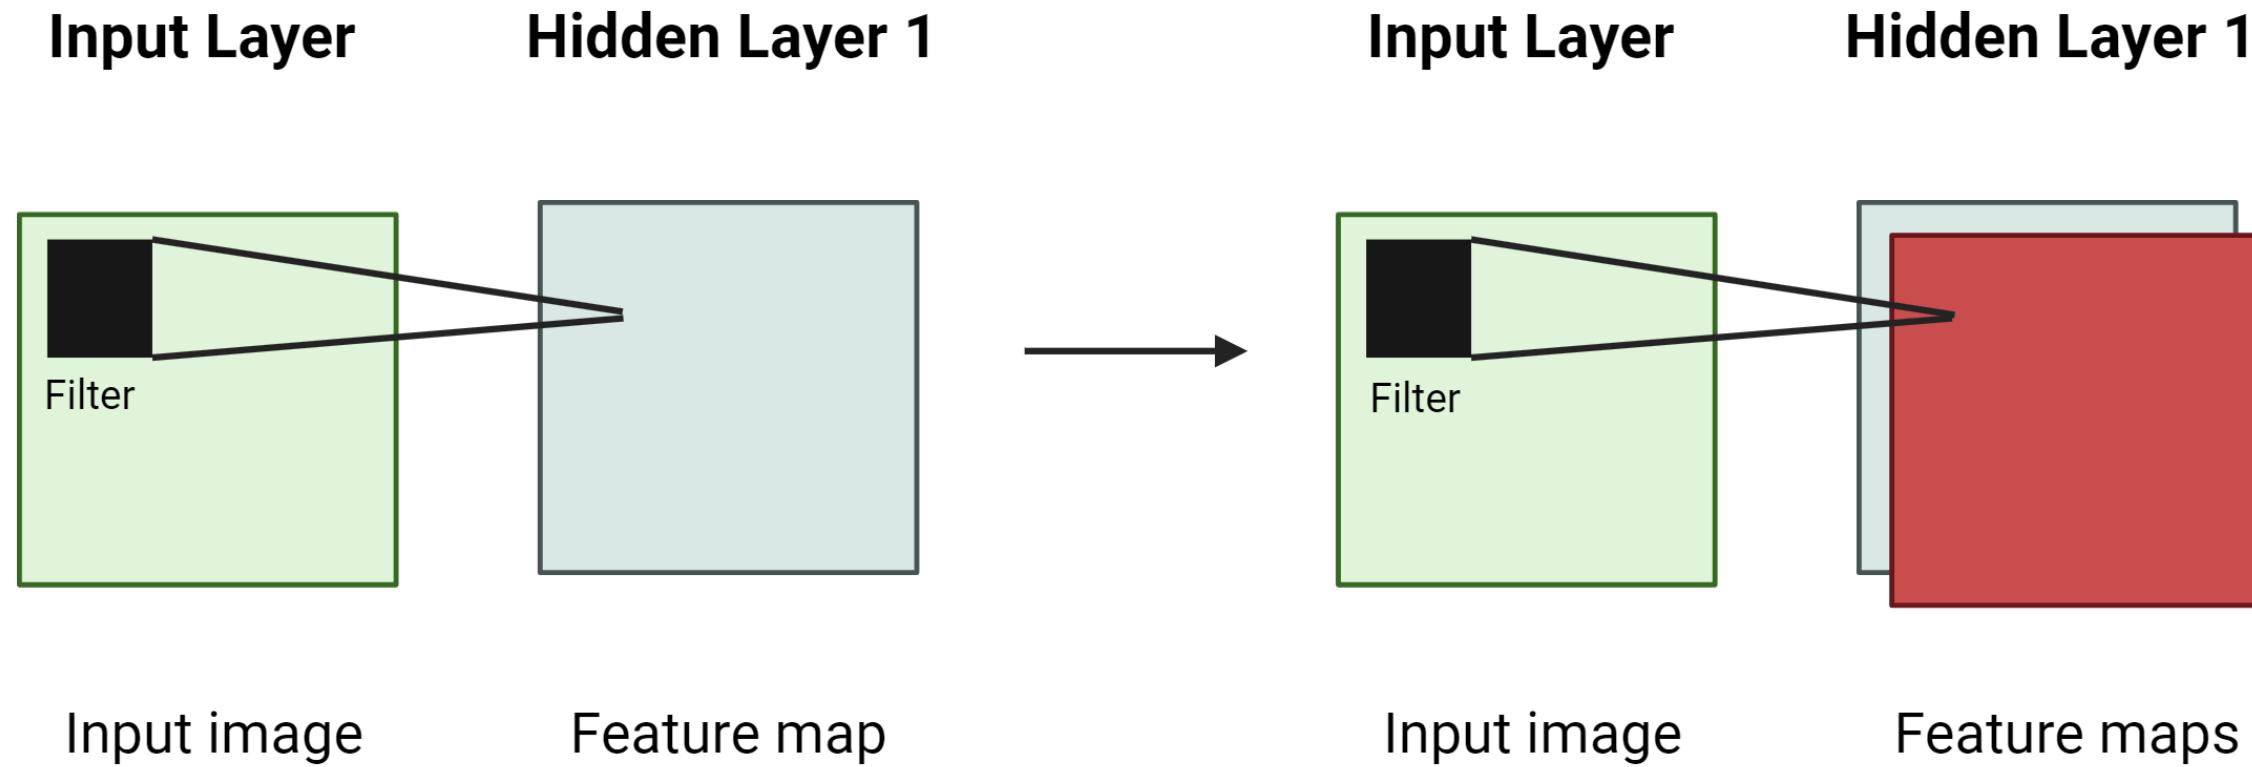
Sharpen	Edge Detection
$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$	$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$
	

# Convolutional Neural Network

- Implicit bias term  $b$  not shown
- Activation function still applied



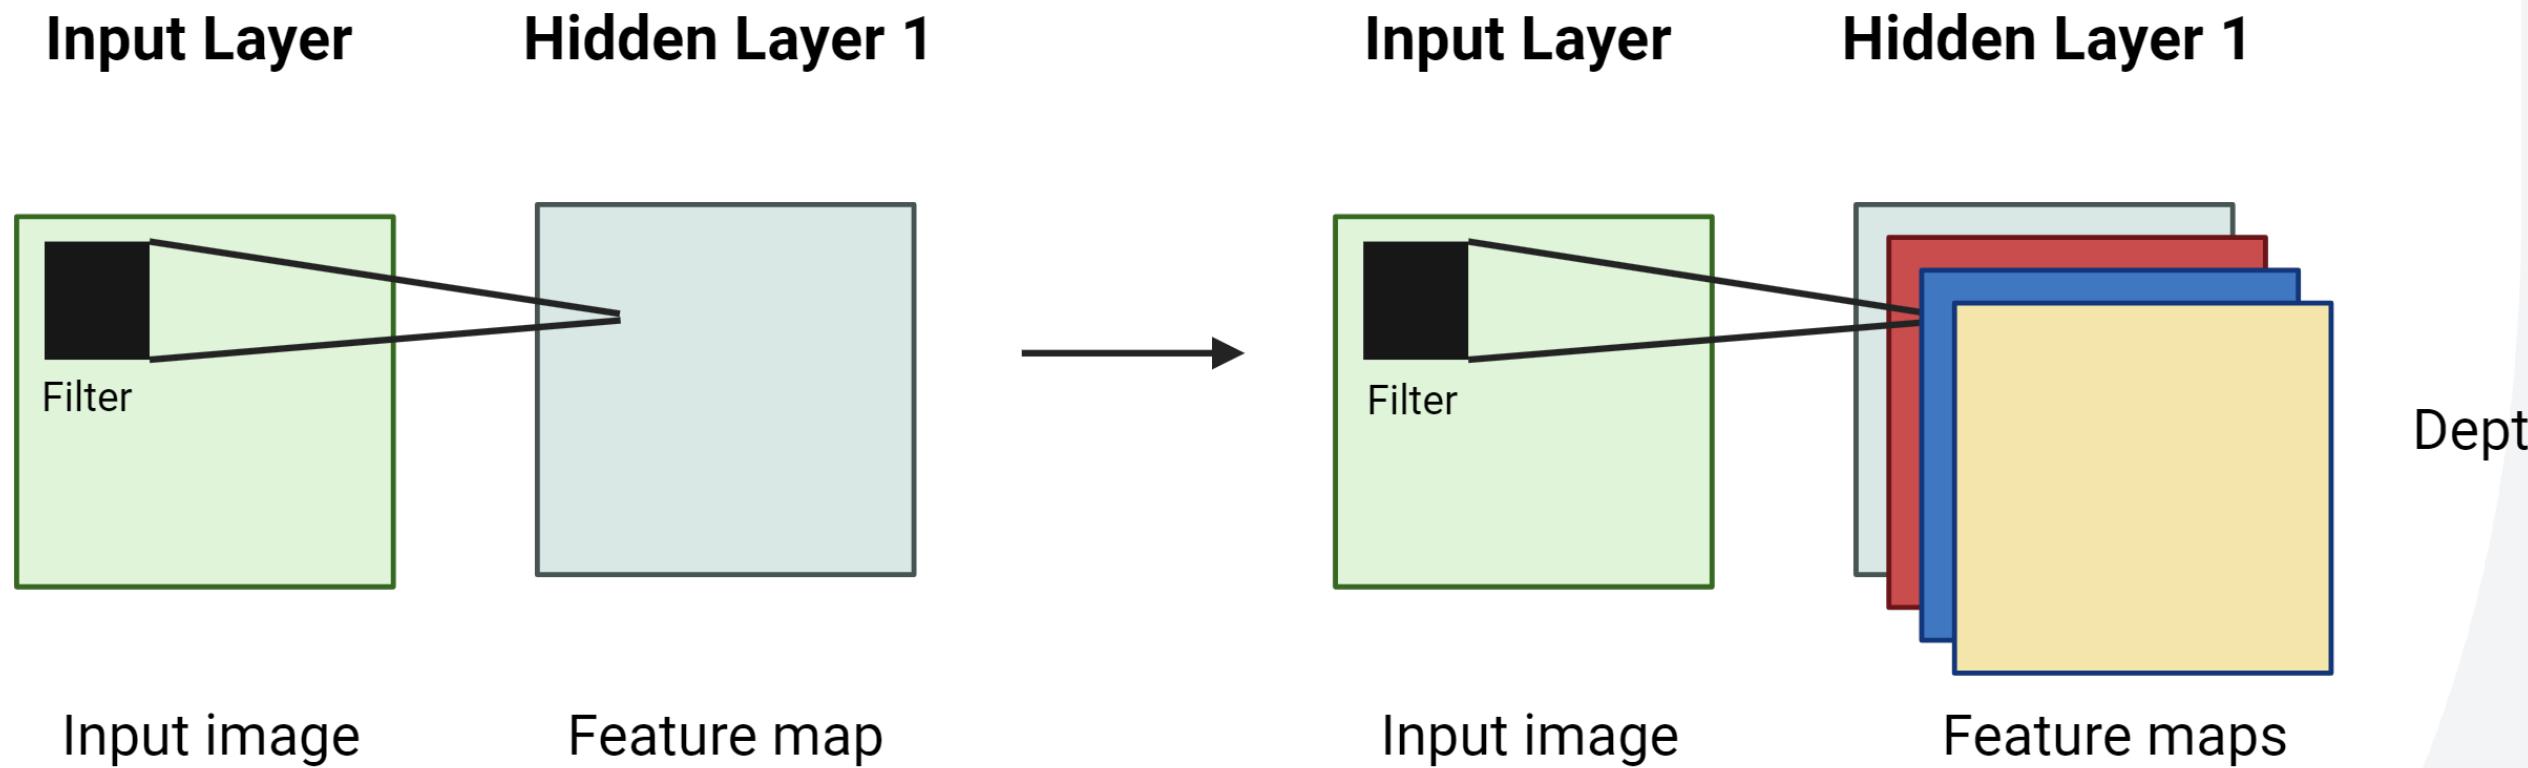
# Convolutional Neural Network



*Convolution operation performed by sliding same filter/kernel over image to produce output feature map*

*Different convolution filters/kernels applied to generate multiple different feature maps*

# Convolutional Neural Network



*Convolution operation performed by sliding same filter/kernel over image to produce output feature map*

*Different convolution filters/kernels applied to generate multiple different feature maps*

Values in each hidden layer are analogous to values we would have for each node in a neural network, which are used in next layer to do calculations

# Convolutional Neural Network

The input might not be scalars but instead RGB values or a 3D image (a stack of ‘channels’). In this case, we treat it as a layer with a depth of 3 (e.g., filter dimensions 3x3x3).

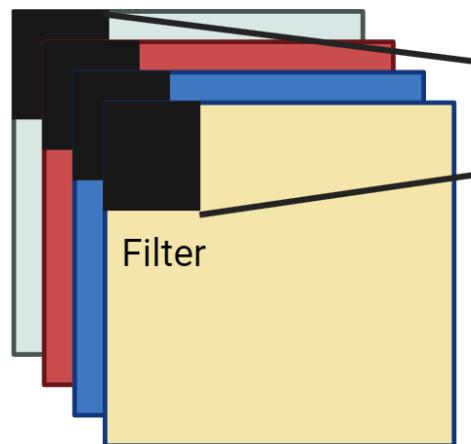
**Input Layer**



Input image

*Filters slides over input image to create feature maps.*

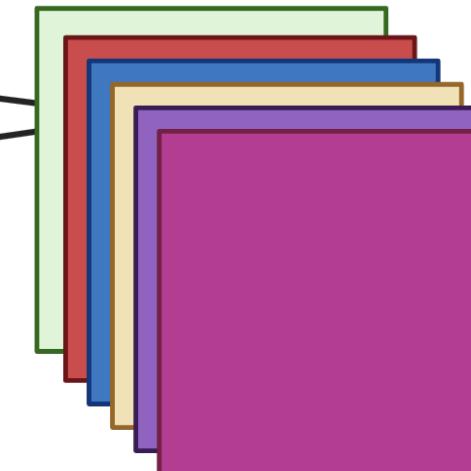
**Hidden Layer 1**



Feature maps

*In hidden layer 1, layers have extra dimension (depth) equal to number of feature maps.*

**Hidden Layer 2**



**Convolution applied to hidden layers:** Each filter applied across all depths. Each filter results in a new feature map in next layer.

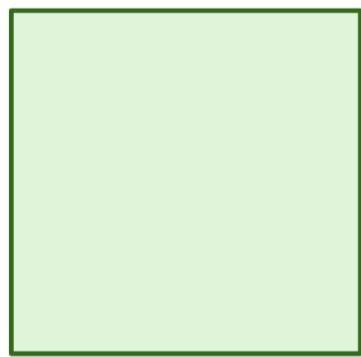
- Depth 4 in hidden layer 1  
→ 4 filters applied
- Depth 6 in hidden layer 2  
→ 6 filters applied

Once we get to the next layer (hidden layer 2), we are applying convolution to multiple feature maps from the previous layer. This allows the network to learn more complex features by combining information from different feature maps.

# Convolutional Neural Network

The input might not be scalars but instead RGB values or a 3D image (a stack of ‘channels’). In this case, we treat it as a layer with a depth of 3 (e.g., filter dimensions 3x3x3).

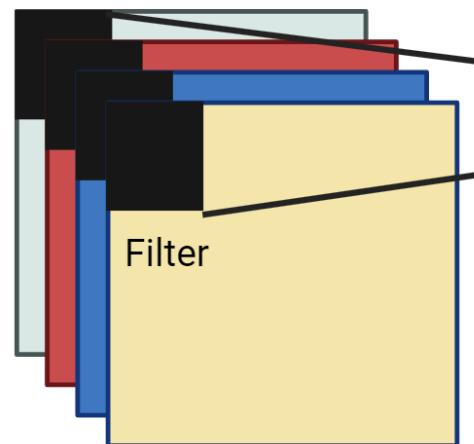
**Input Layer**



Input image

*Filters slides over input image to create feature maps.*

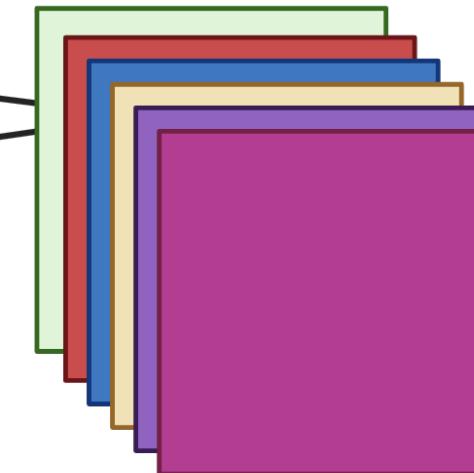
**Hidden Layer 1**



Feature maps

*In hidden layer 1, layers have extra dimension (depth) equal to number of feature maps.*

**Hidden Layer 2**



*In hidden layer 2, convolutions are applied across all feature maps from hidden layer 1, producing new feature maps that combine information from all previous maps.*

**Convolution applied to hidden layers:** Each filter applied across all depths. Each filter results in a new feature map in next layer.

- Depth 4 in hidden layer 1  
→ 4 filters applied
- Depth 6 in hidden layer 2  
→ 6 filters applied

Once we get to the next layer (hidden layer 2), we are applying convolution to multiple feature maps from the previous layer. This allows the network to learn more complex features by combining information from different feature maps.

# Pooling + Parameters

# Pooling Layer

Pooling layers (down-sample layers) perform a down-sampling operation, reducing the spatial dimensions (width and height) of the input feature map while retaining the important features.

## Max pooling layer

Original image

4	8	3	5
2	4	5	3
1	9	1	7
7	6	0	0

Output

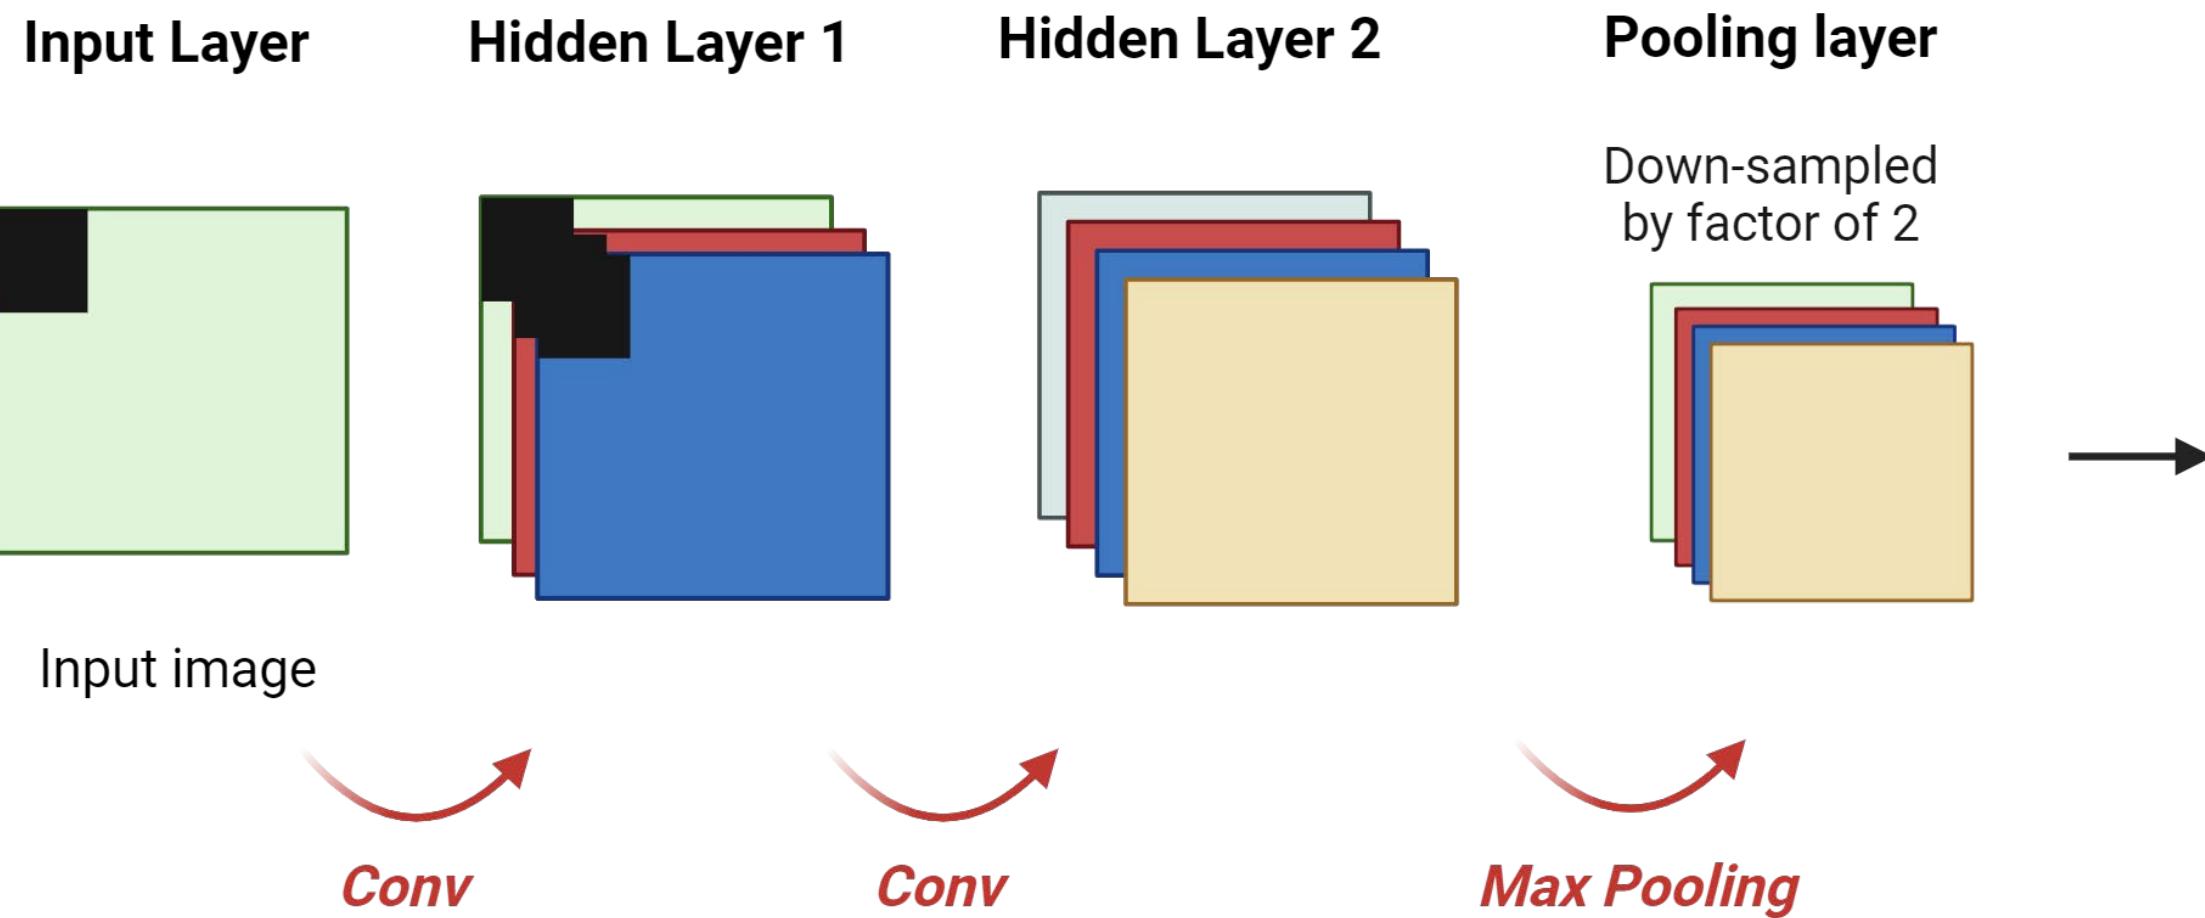
8	5
9	7

Take maximum value within a 2x2 filter  
Move in steps of 2 (**stride**)  
Down samples image by factor of 2

- Useful for reducing the number of parameters required in a CNN.
- Pooling layers have no parameters.
  - The values in the filters (or kernels) for convolution are parameters that are learned during training.
  - These are updated through optimisation/backpropagation, similar to weights in traditional neural networks.
  - Different feature maps are generated during optimisation due to updates in the filters.
- Pooling doesn't have to be max pooling; average pooling is also an option.

# Convolutional Neural Network

*Each layer only connects to next layer*



Convolutional layers are often interspersed with max pooling layers

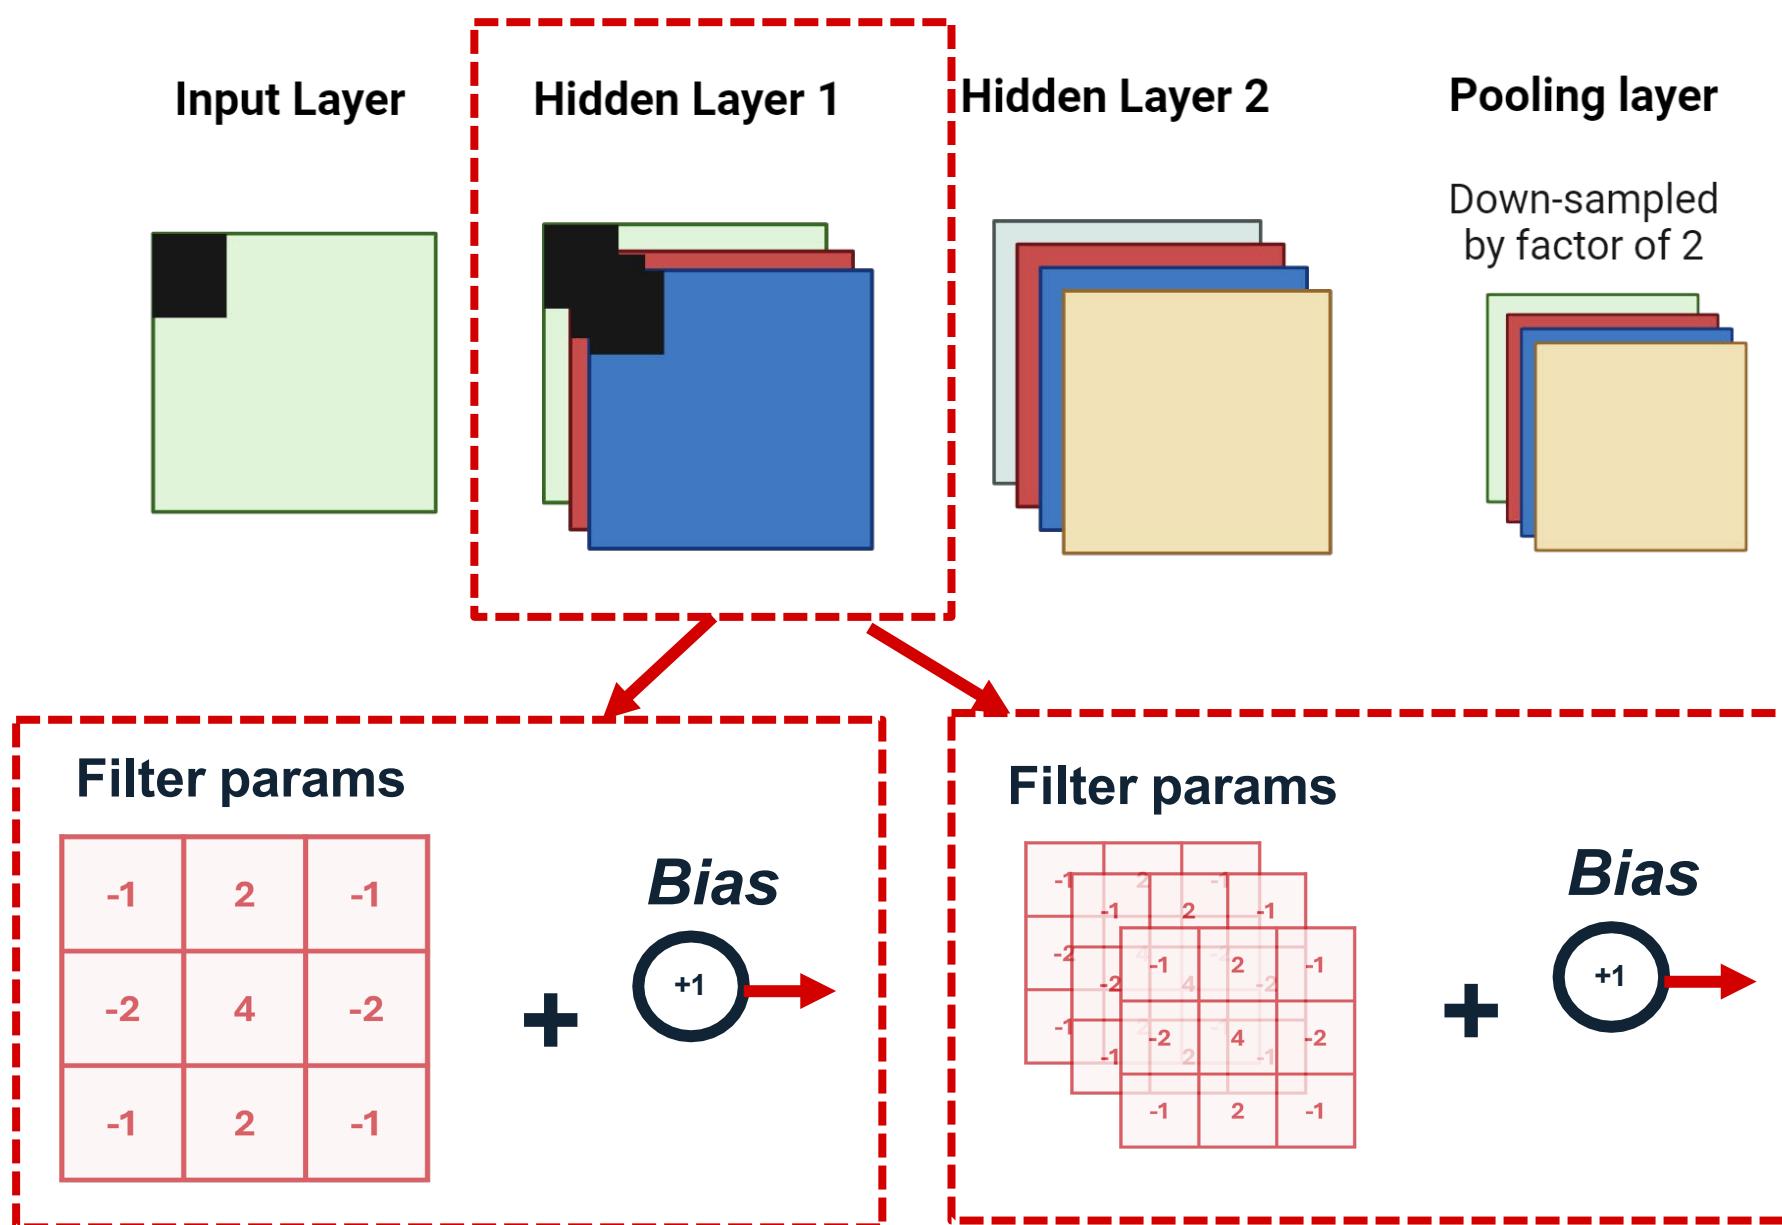


THE UNIVERSITY  
of ADELAIDE

150  
YEARS

# Parameters

- Values in the filters/kernels for convolution = parameters learnt during training.
- Updated through optimisation/backpropagation, generating different feature maps



**For a  $3 \times 3$  filter**

**Scalar input image (depth = 1)**

# Params = Filter Size + Bias

$$3 \times 3 \times 1 + 1 = 10$$

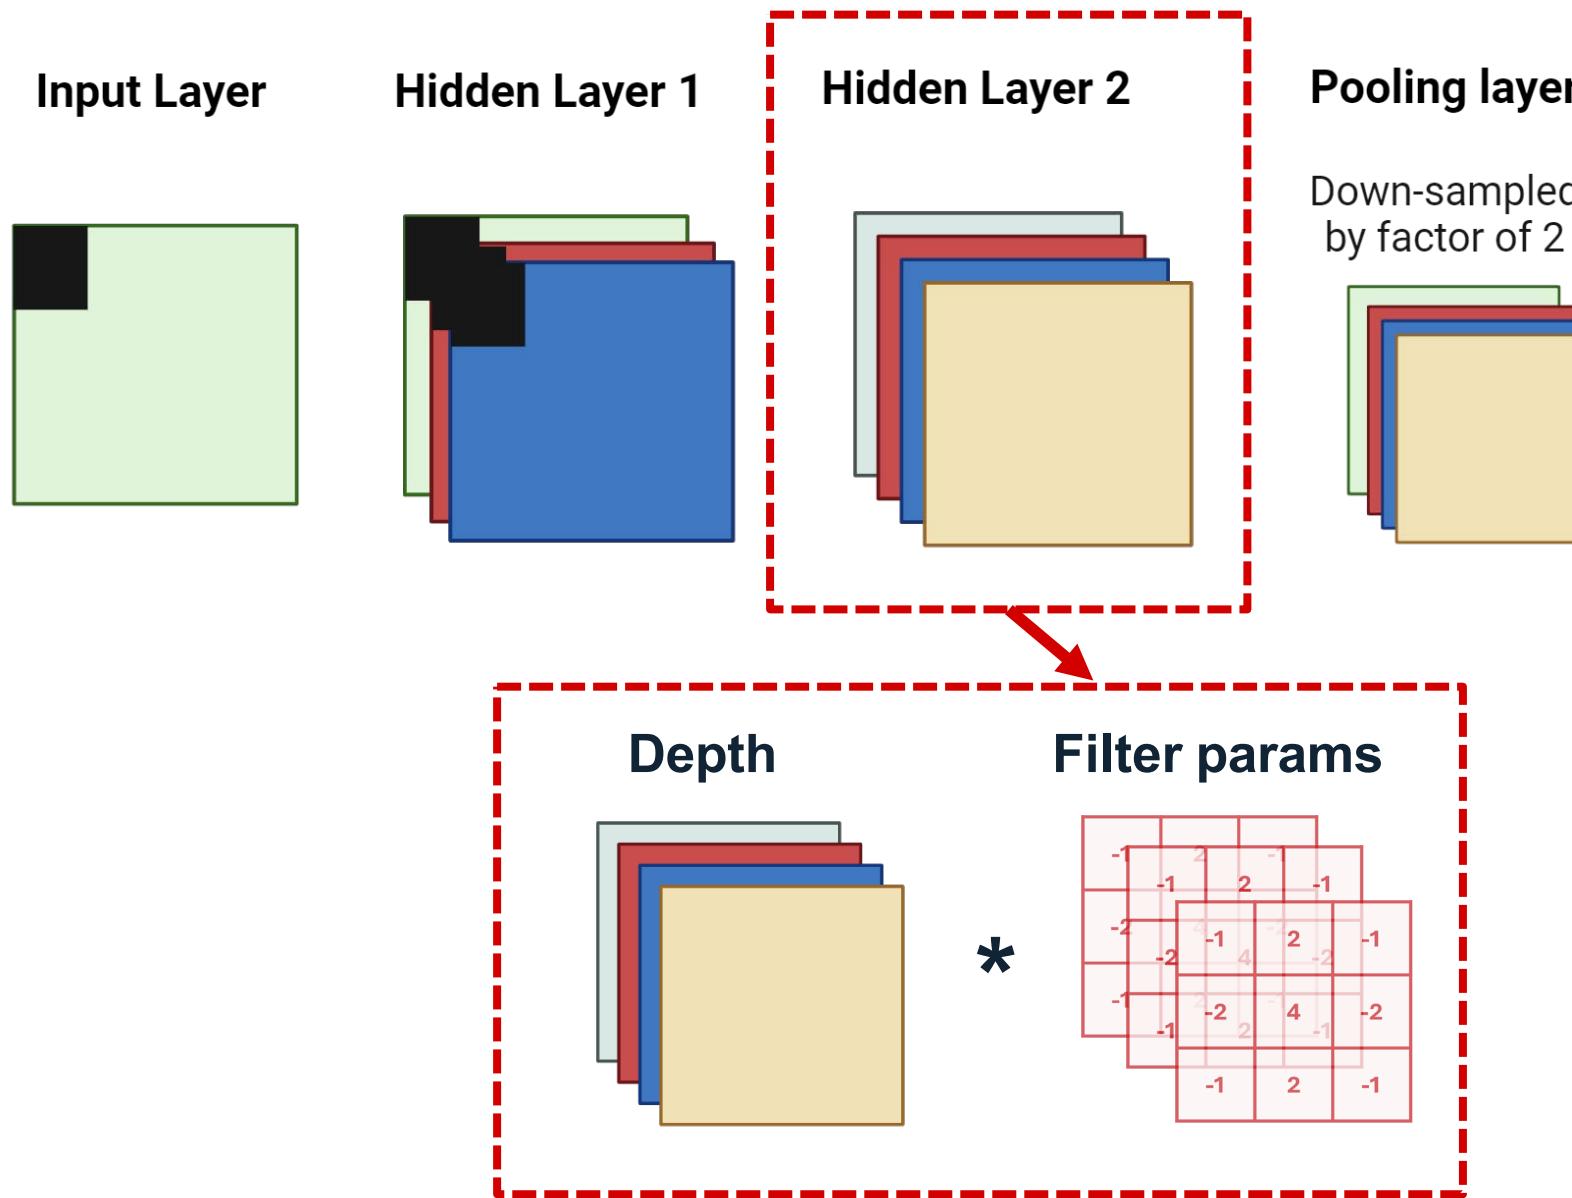
**Hidden layer 1**

# Parameters in each filter:

$$(3 \times 3 \times 3 + 1) = 28$$

# Parameters

- Values in the filters/kernels for convolution = parameters learnt during training.
- Updated through optimisation/backpropagation. Different feature maps are generated during optimisation due to updates in the filters.



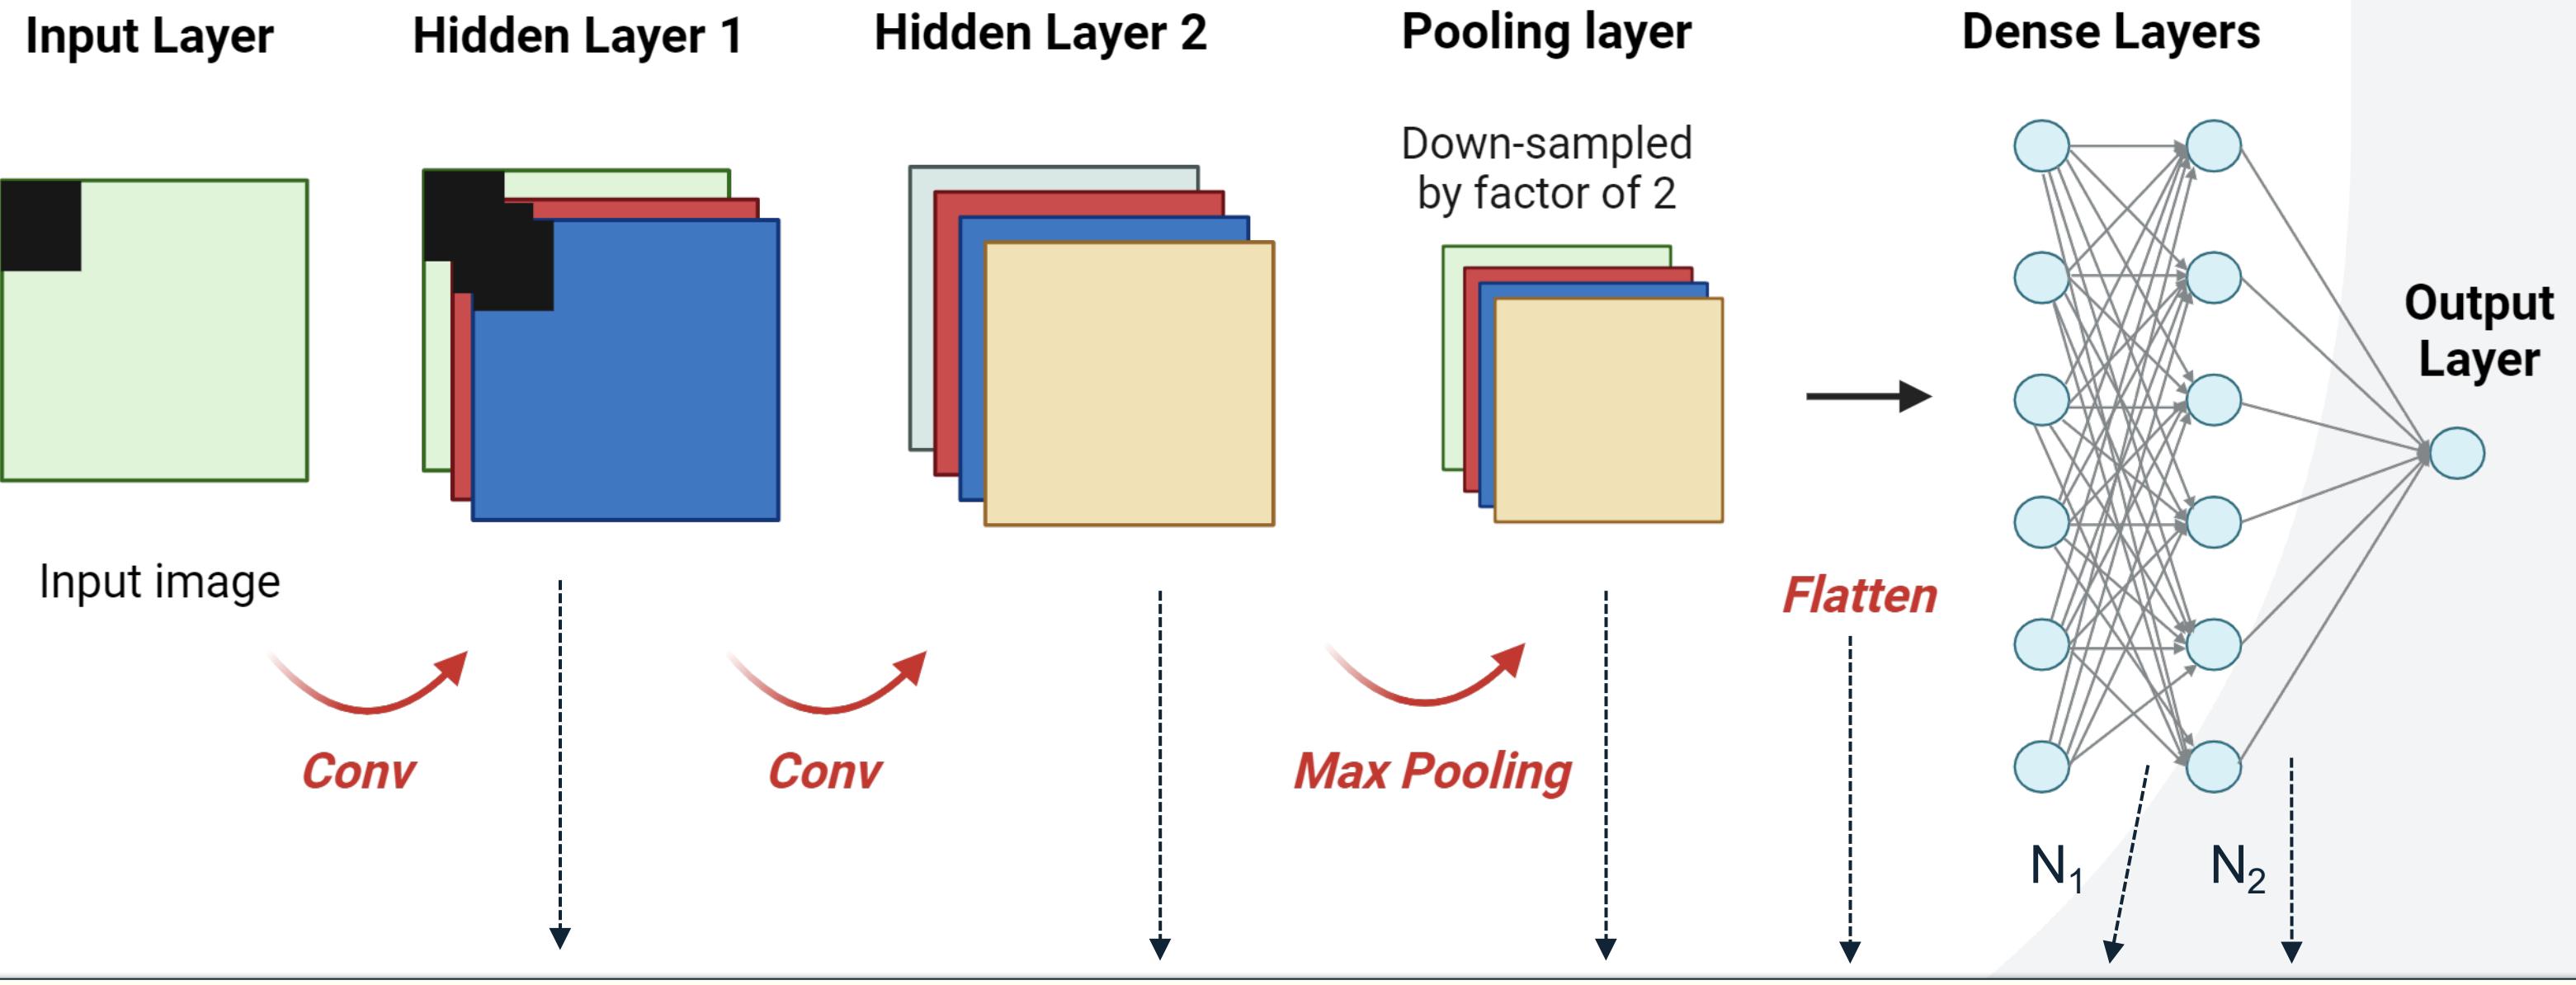
**For a 3x3 filter in hidden layers**

**Hidden layer 2**

# filters/depth \* #params in each filter:  
 $(3*3*3+1)*4 = 28*4 = 112$

# Convolutional Neural Network

*Each layer only connects to next layer*



Total # params:  $(3*3*3+1) = 28$

$(3*3*3+1)*4 = 112$

0

0

$(N_1+1)*N_2$

$N_2+1$

# Which answer is correct?

What is the main purpose of an activation function in a neural network?

- To initialise weights
- To introduce non-linearity
- To normalise input data
- To optimise the loss function

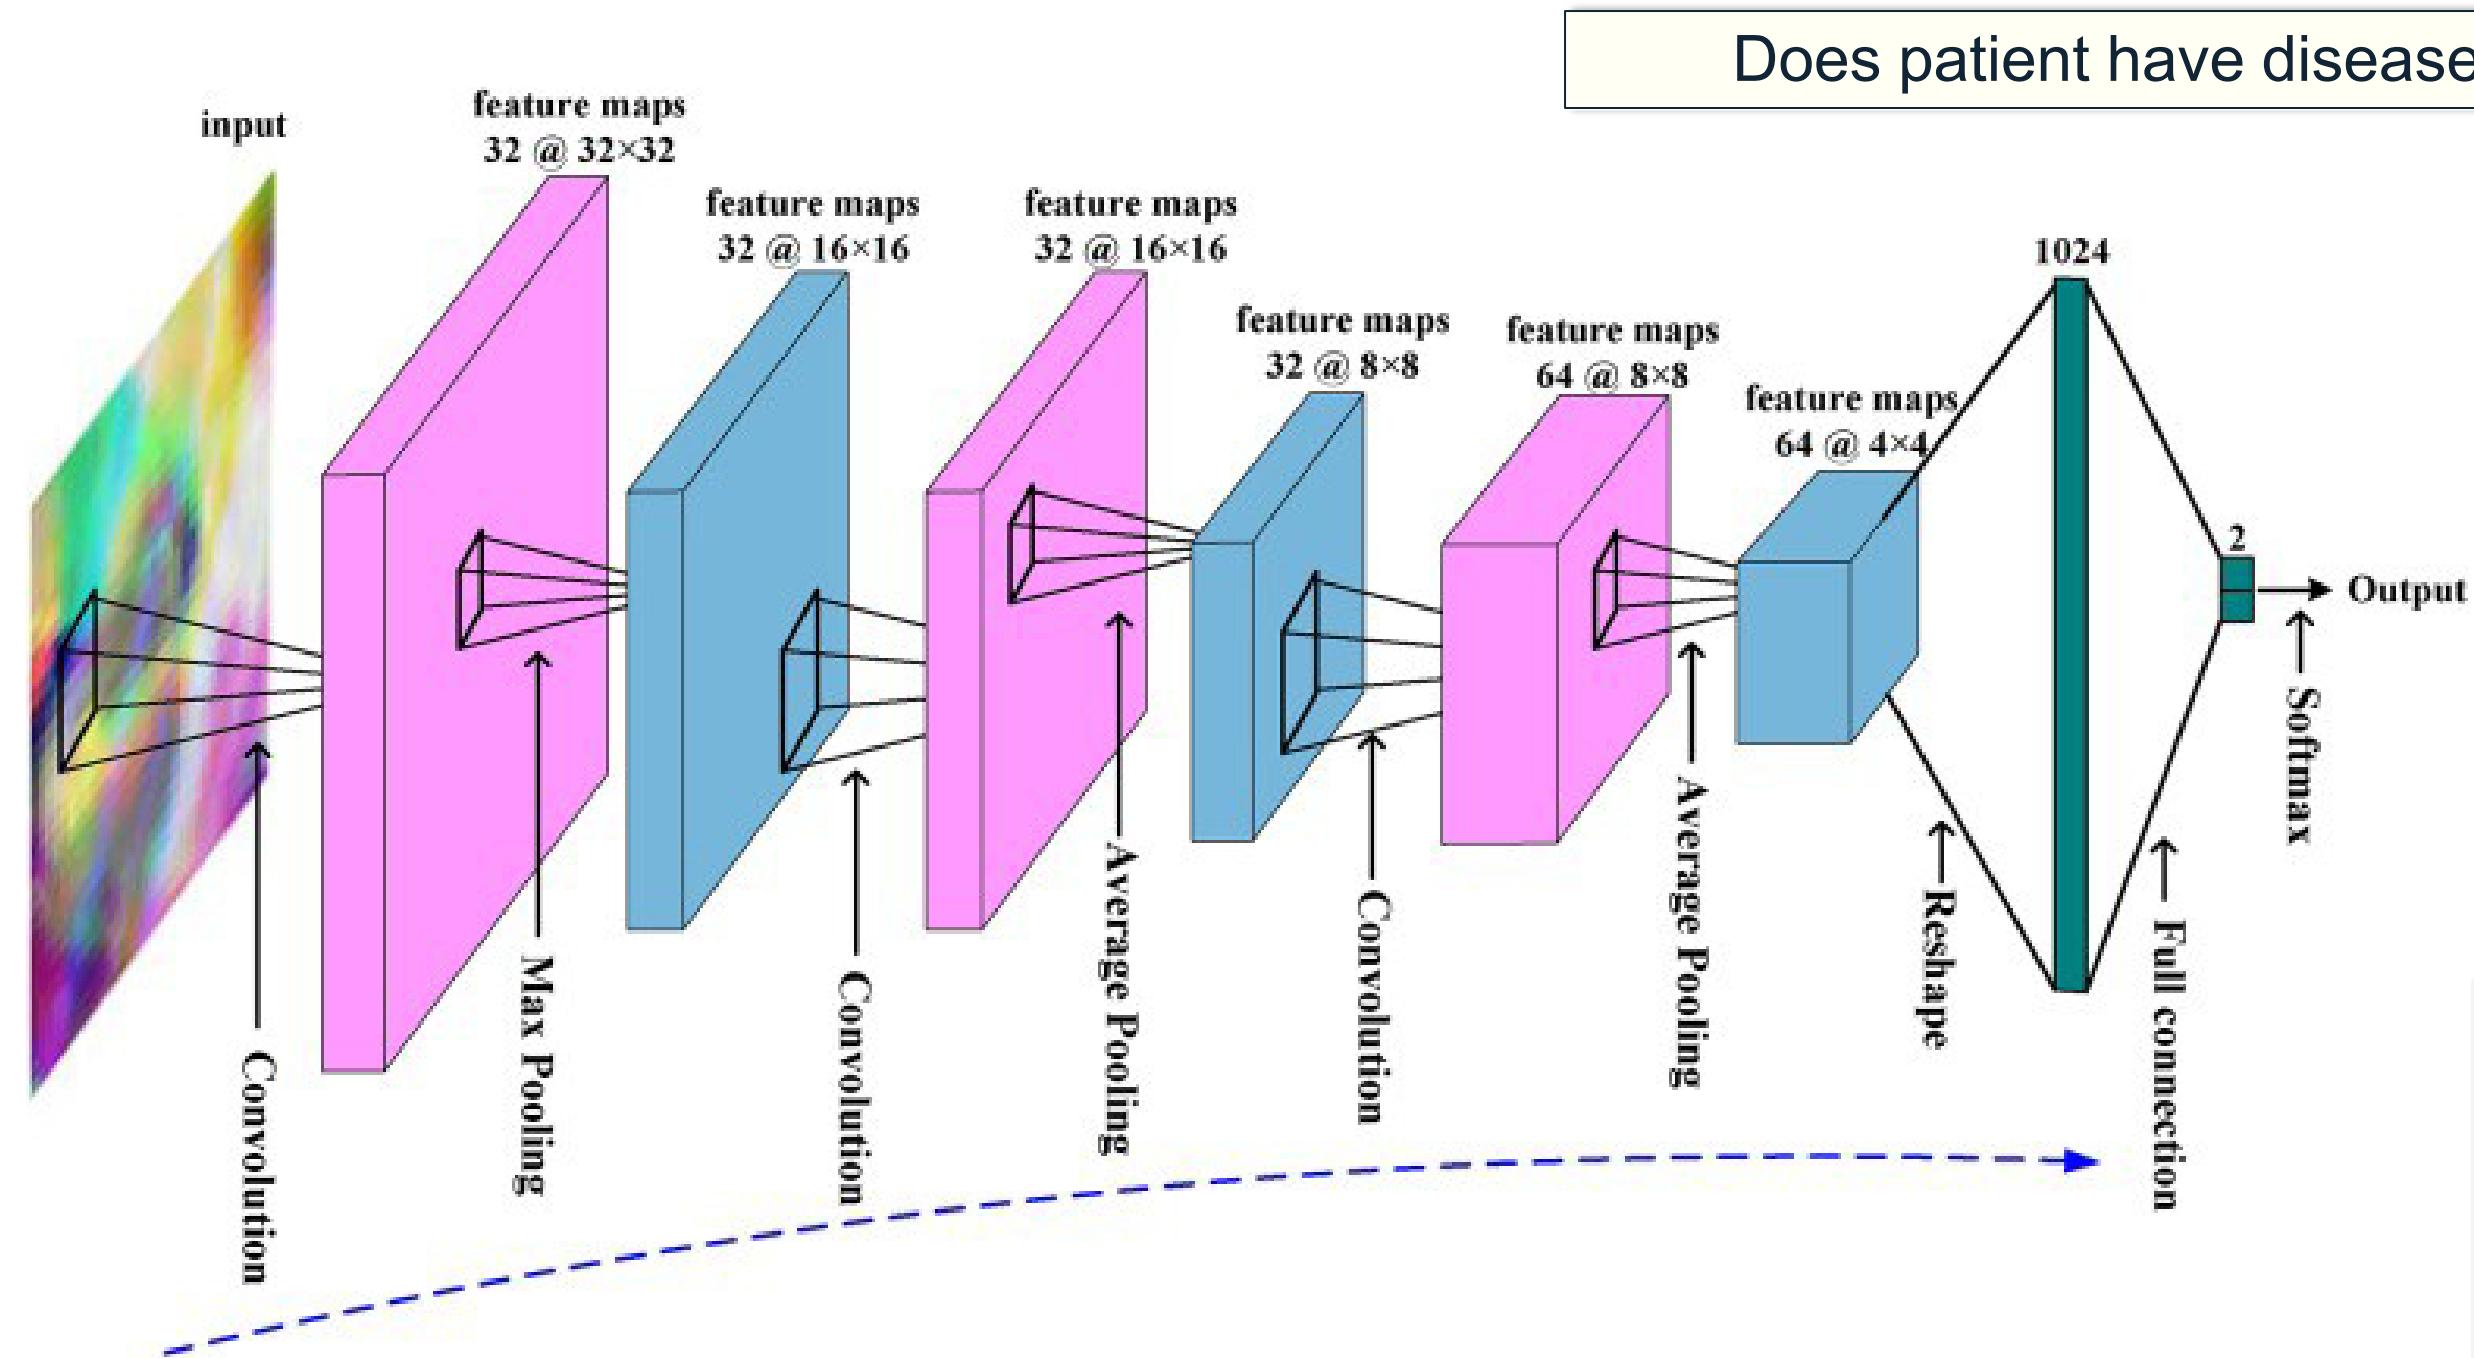


THE UNIVERSITY  
of ADELAIDE

15C  
YEARS

# Applications + Considerations

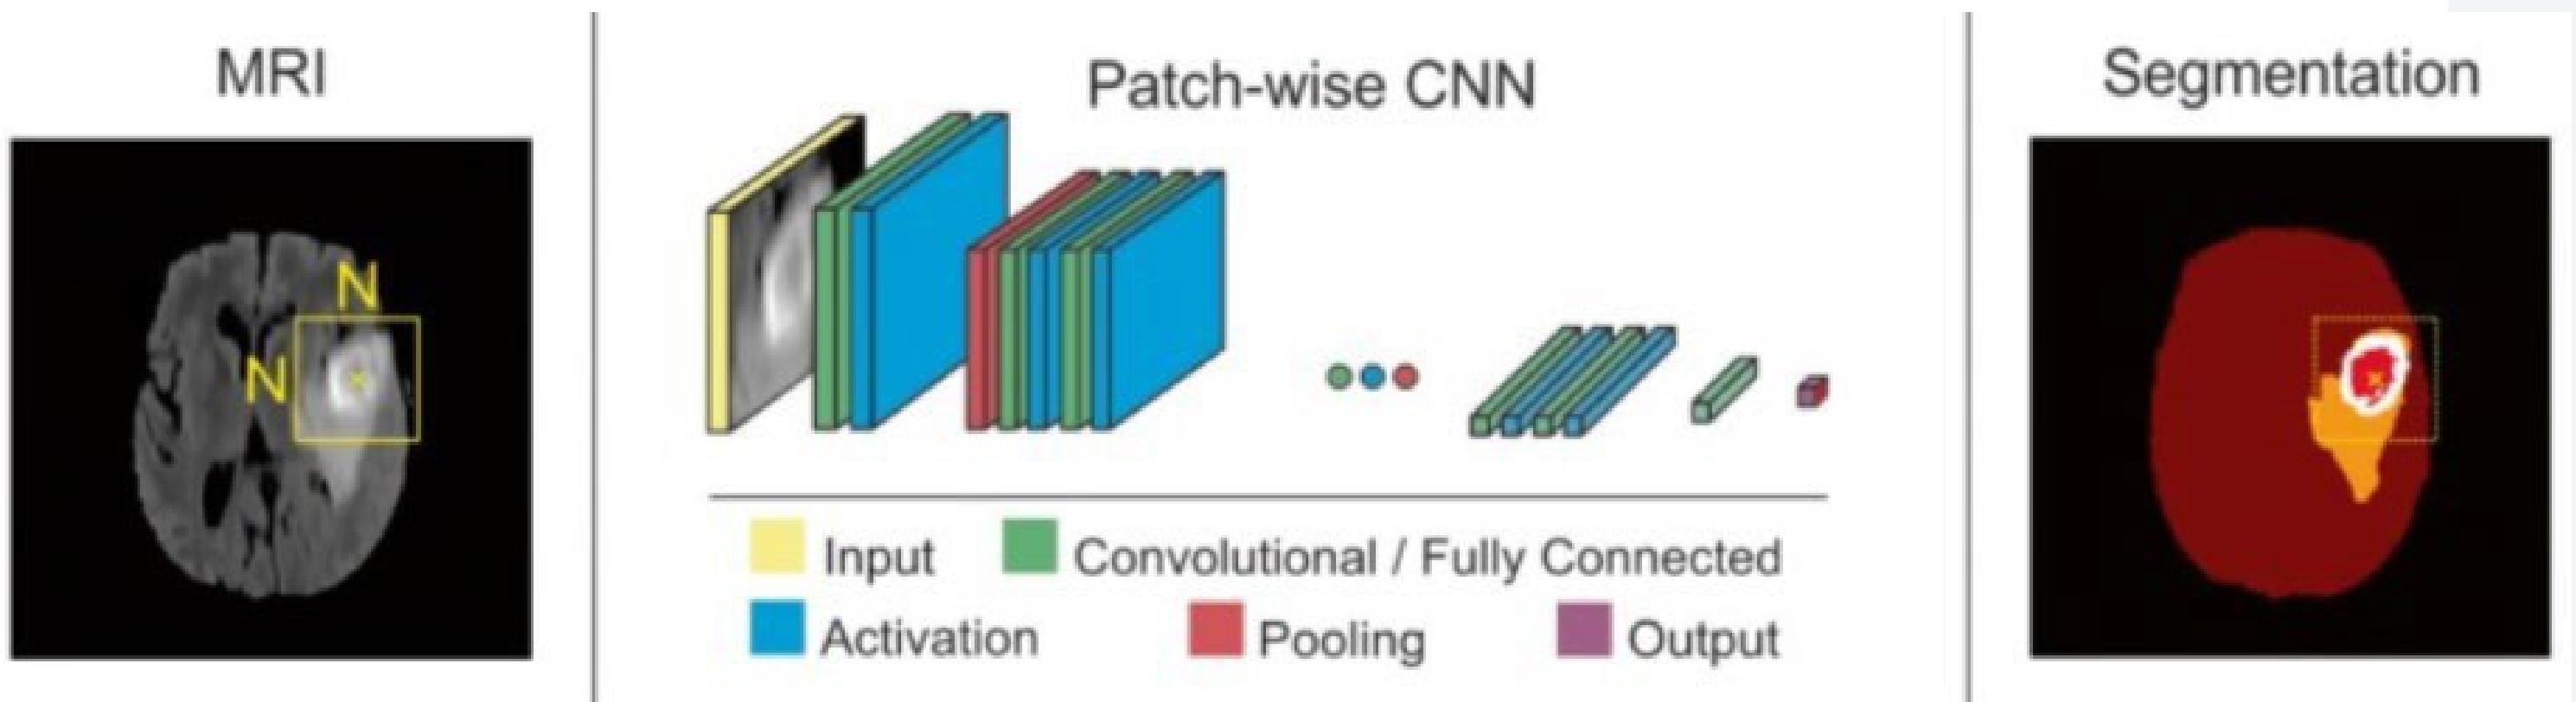
# Example: Classification



Lin et al., Convolutional Neural Networks-Based MRI Image Analysis for the Alzheimer's Disease Prediction From Mild Cognitive Impairment, *Front. Neurosci.* 2018

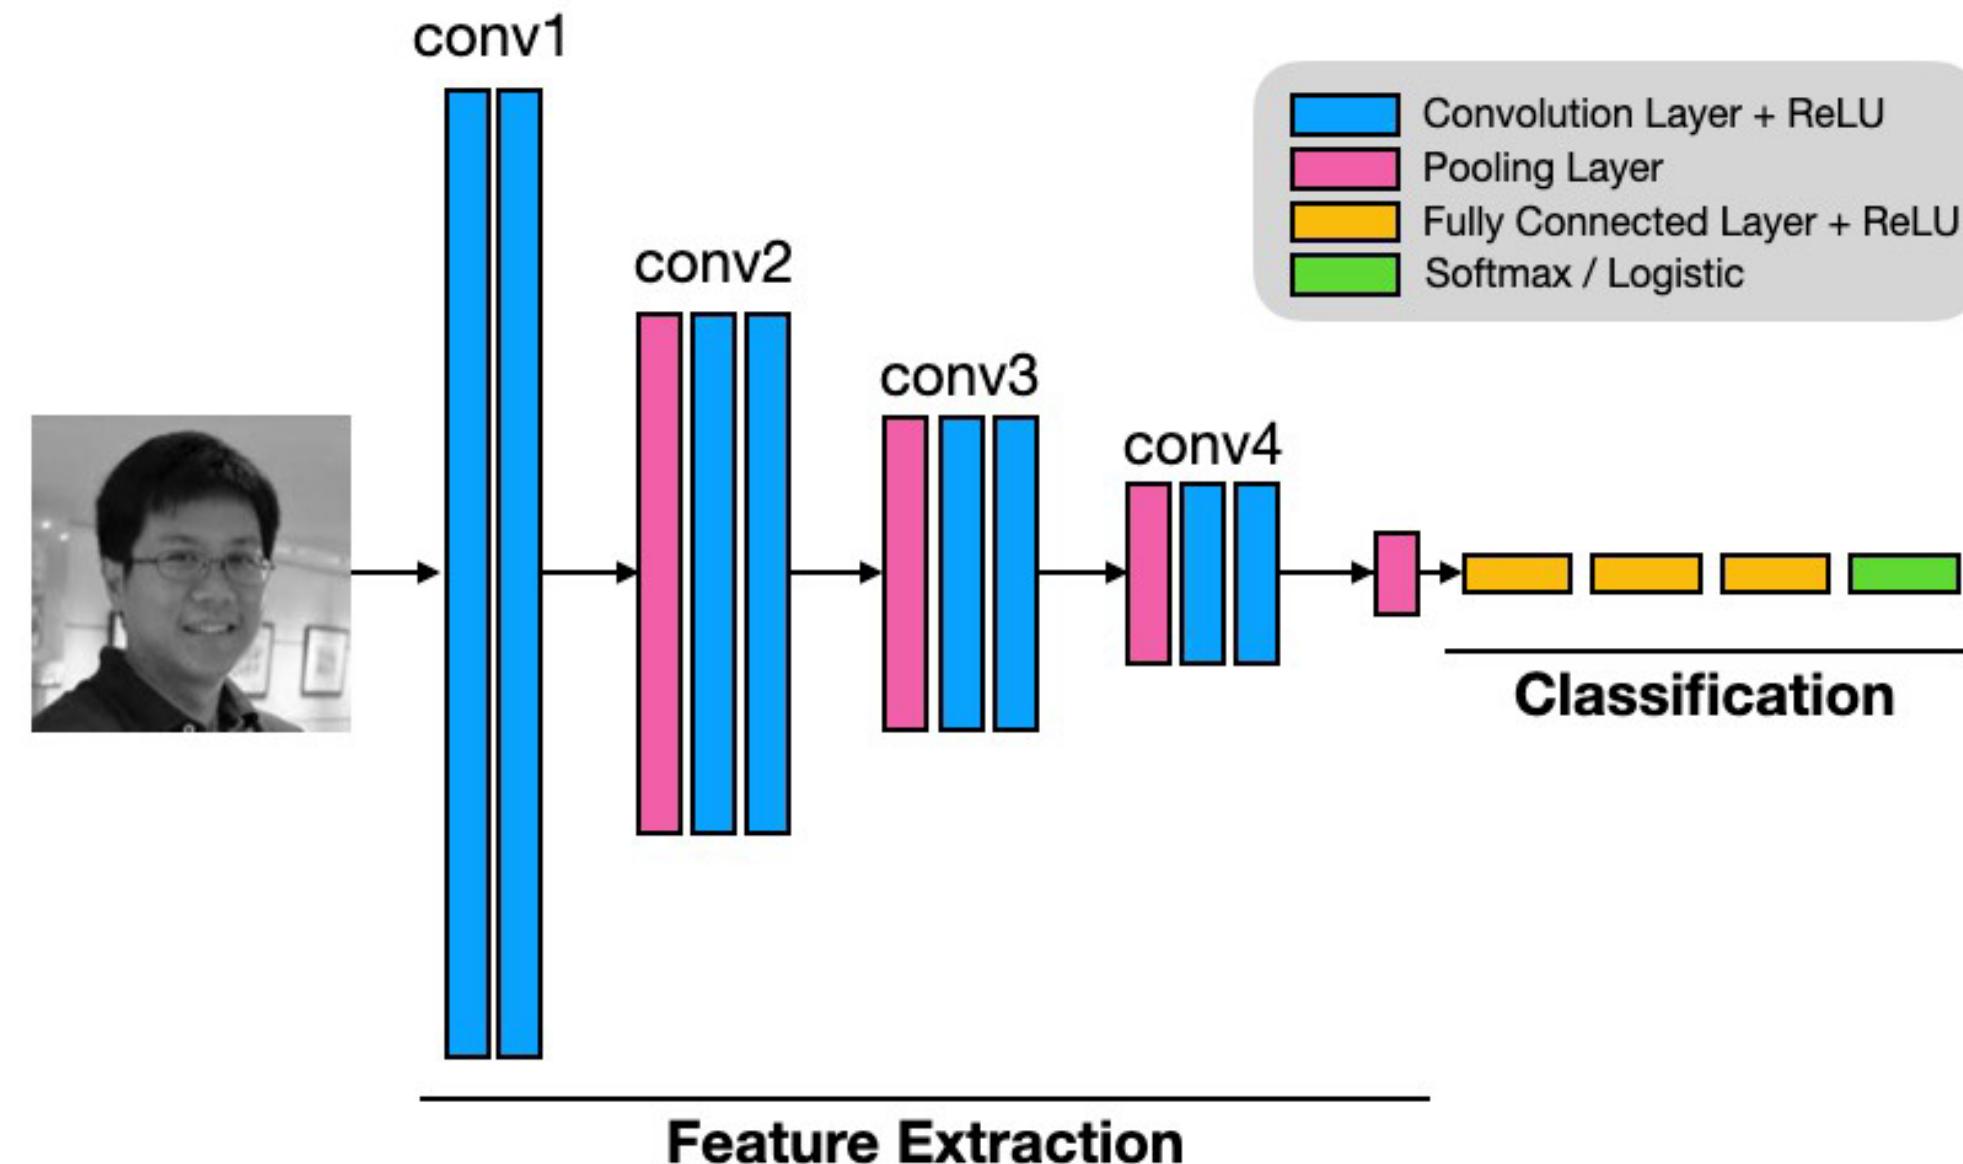
# Example: Segmentation

- What type of brain tissue (coloured) is present at each point/location in image?
  - Like a point-by-point classification
- Patch-based application, meaning this methodology applied across different patches/chunks ( $N \times N$ ) of the image at a time and give label for each centre point

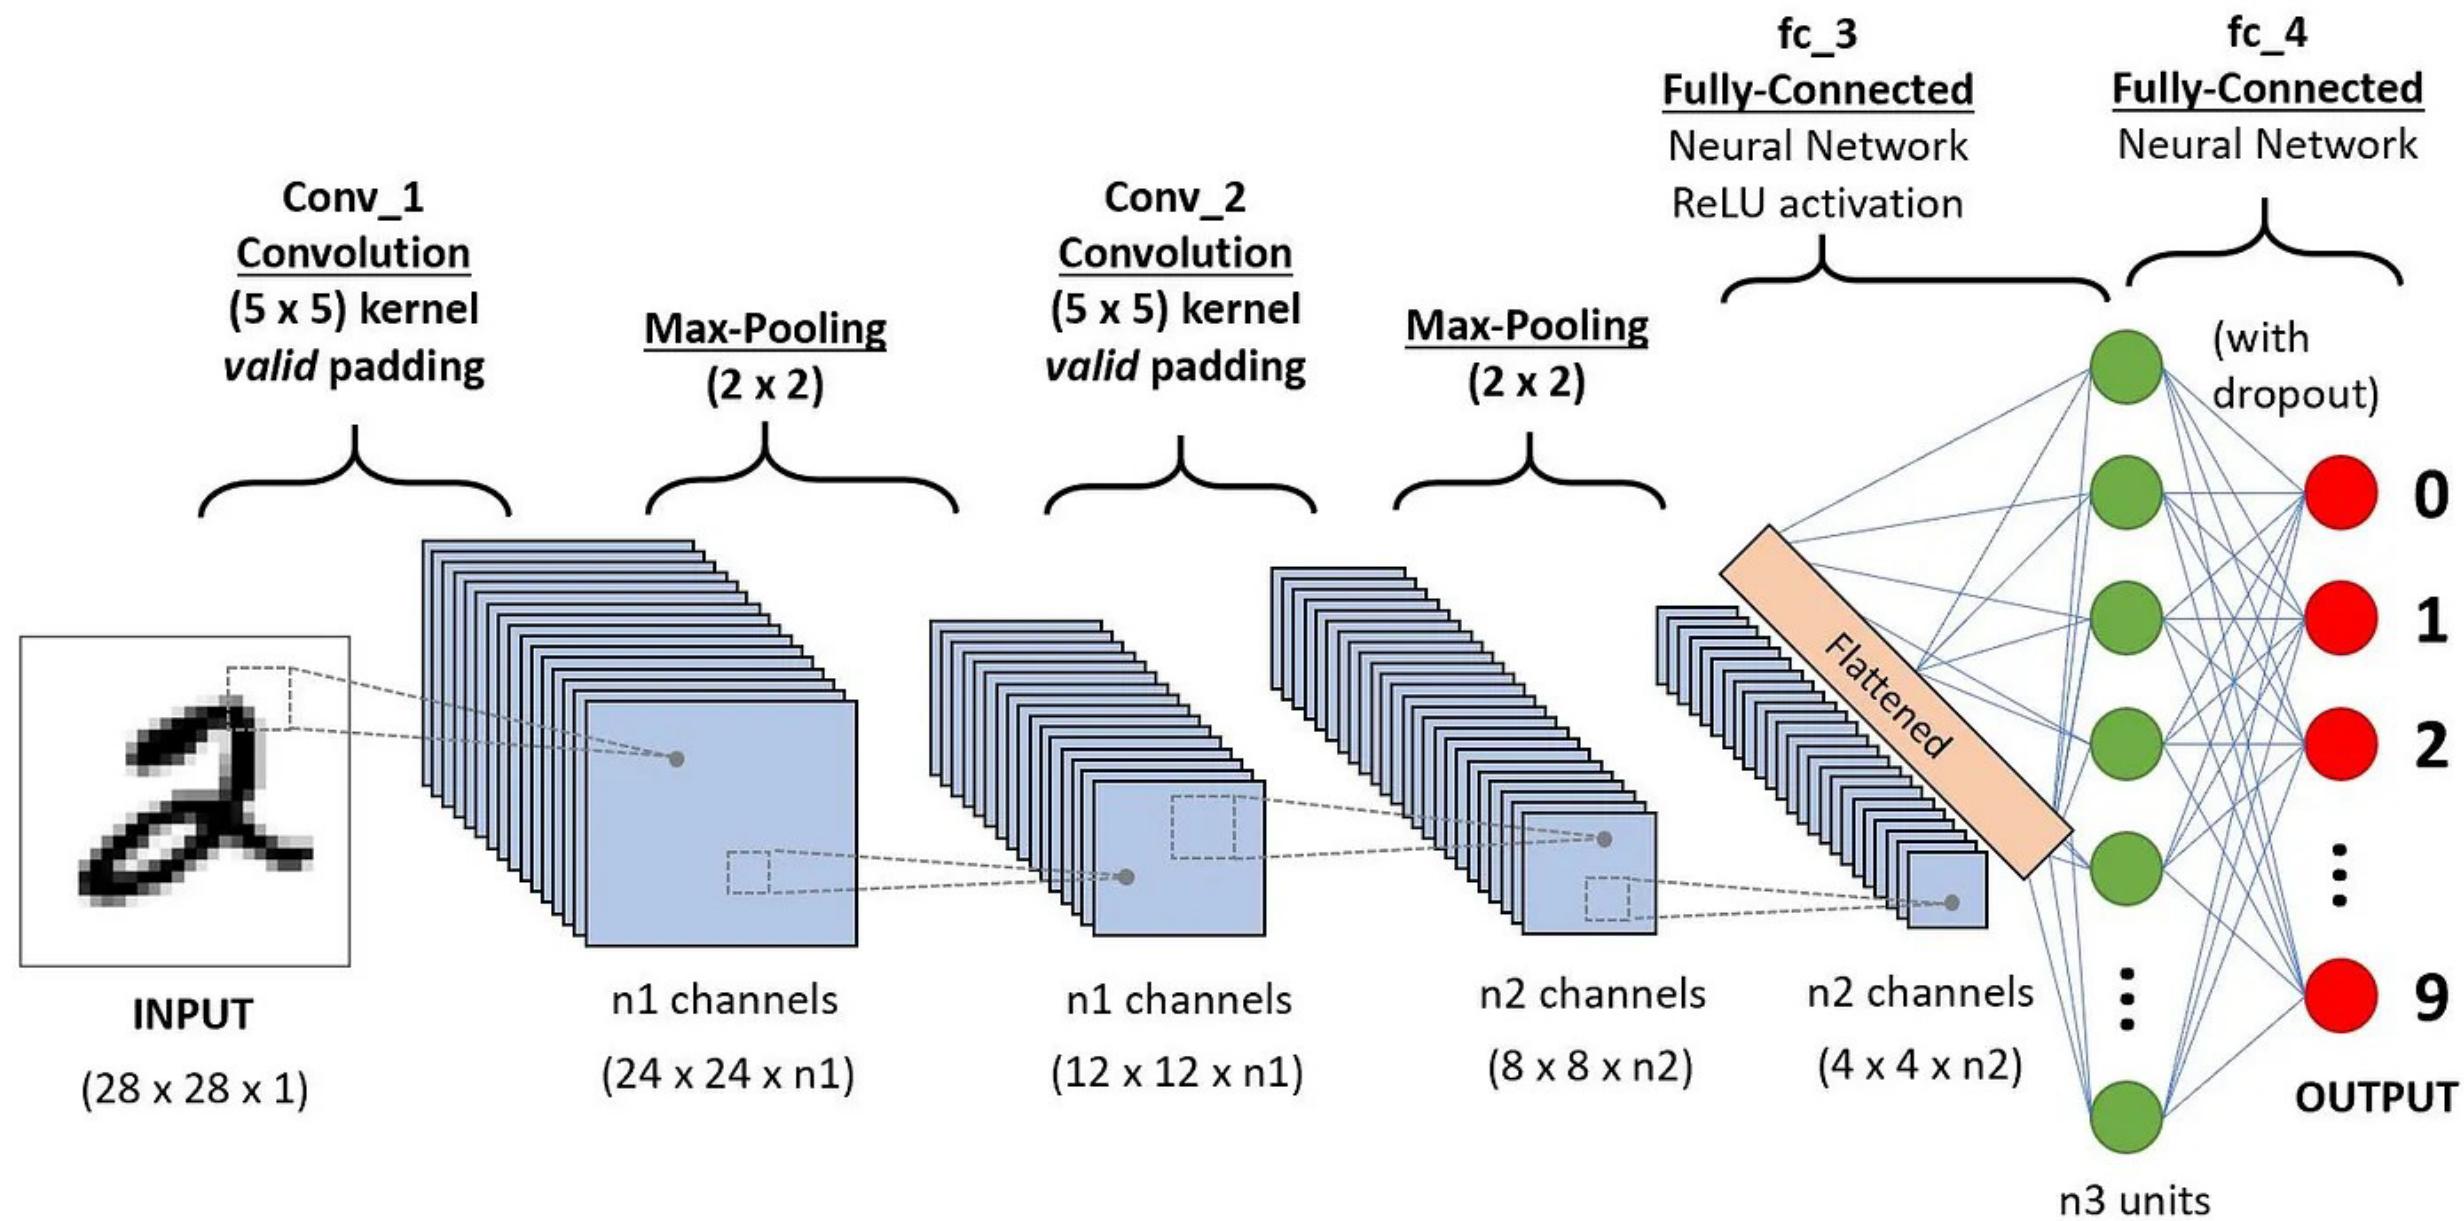


# Example: Facial Recognition

- Extracts unique facial features from image and decides if it belongs to person A, B or C.

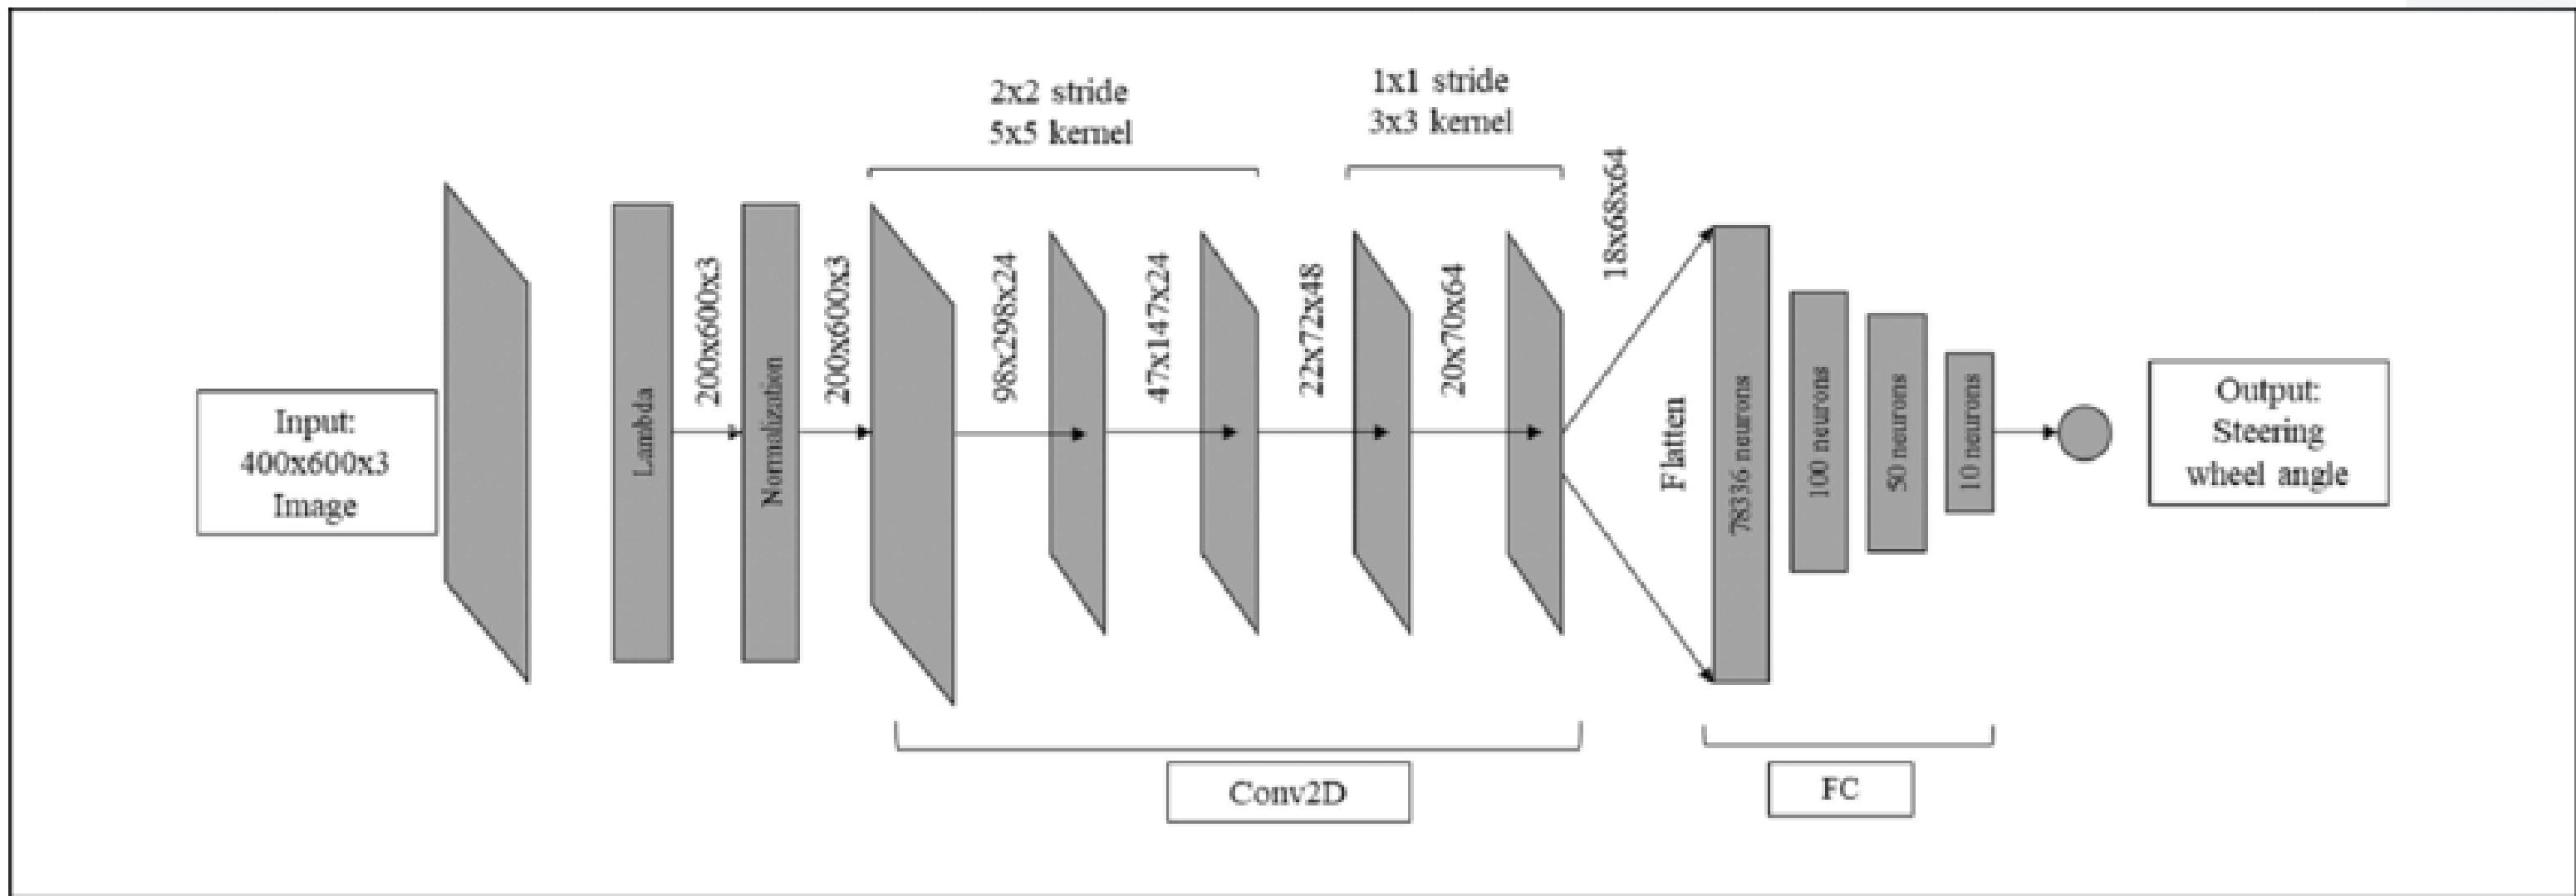


# Example: Text Recognition



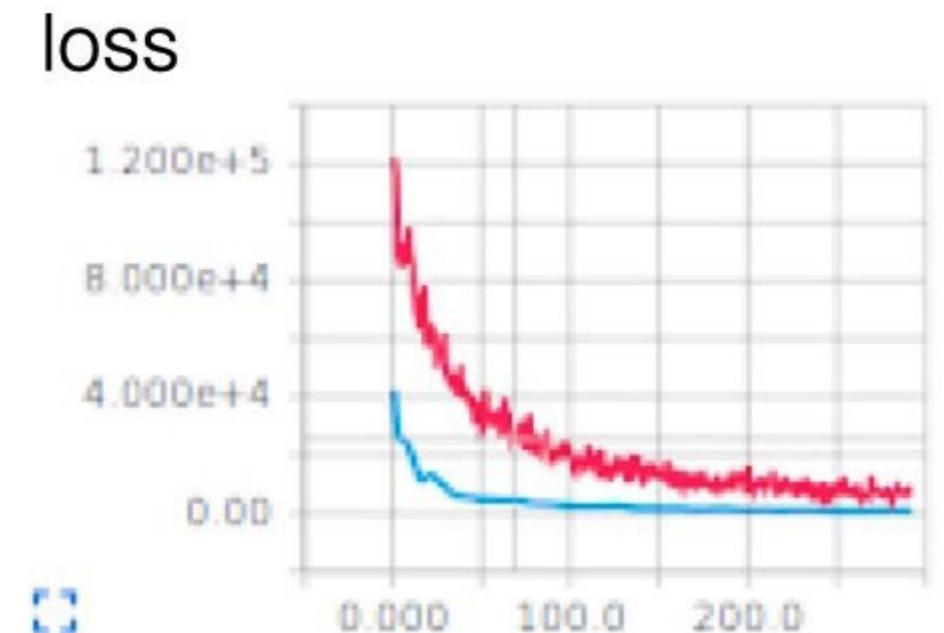
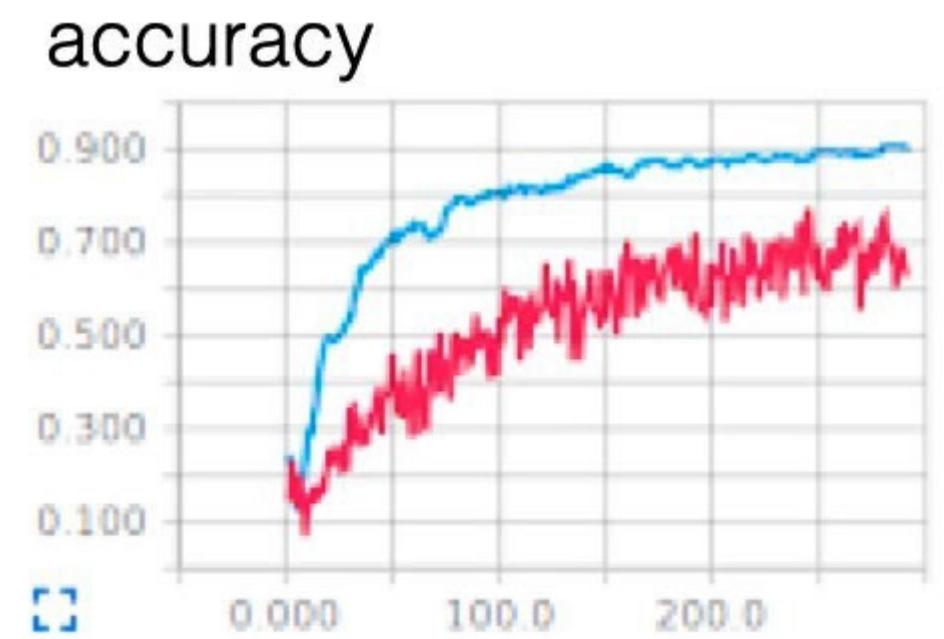
# Example: Autonomous Vehicles

Uses CNN to detect hazards on the road and produces a segmented image to assist the vehicle in navigating safely



# Convergence and Overfitting

- Batch size, epochs and convergence are the same as discussed before
- Overfitting more likely when number of parameters exceeds number of training data inputs
  - Validation acc/loss much worse than training acc/loss
- Underfitting occurs when insufficient number of parameters or insufficient richness in data
  - Training acc/loss poor
- Can help guide network choices
- Empirical selection most common



# Ensembles

- Results of optimisation (training) will depend on many things (initialisation, number of epochs, order of data, etc.)
- Most of the time it will end up in a local minima (this might be good enough)
- Can combine together different trained models and average outputs to get an improved model
  - Initialised differently
  - Trained differently
  - Different architectures/parameters
- Well studied in traditional machine learning
  - i.e random forest = ensemble of decision trees (one of the best ML models)
- Approach also common and powerful for deep learning
  - Usually, would ensemble just a few in deep learning

# Which answer is correct?

What is the main purpose of an activation function in a neural network?

- To initialise weights
- To introduce non-linearity
- To normalise input data
- To optimise the loss function



THE UNIVERSITY  
of ADELAIDE

15C  
YEARS

# Summary

- 1. Convolutional Neural Networks**
- 2. CNN Structures**
- 3. Pooling + Parameters**
- 4. Applications + Considerations**



THE UNIVERSITY  
of ADELAIDE

**15C**  
YEARS

**Questions?**

[dhani.dharmaprani@adelaide.edu.au](mailto:dhani.dharmaprani@adelaide.edu.au)